

# Cloud Computing

## MapReduce to Spark

---

Minchen Yu  
SDS@CUHK-SZ  
Fall 2024



香港中文大學(深圳)  
The Chinese University of Hong Kong, Shenzhen



Hadoop is great, but it's really  
waaaaay too low level! (circa 2007)

# What's the solution?

Design a higher-level language

Write a compiler

# Hadoop is great, but it's really waaaay too low level!



What we really need is SQL



What we really need is  
scripting language

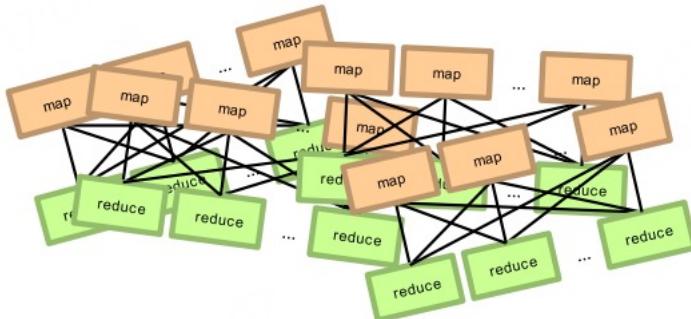




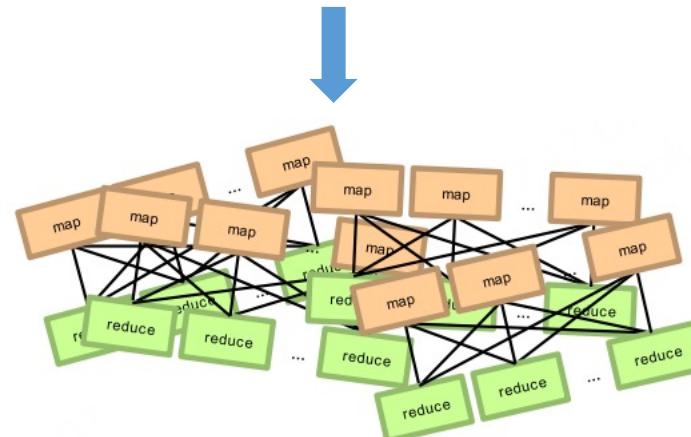
SQL



Why not just use a  
database?



Pig Scripts





Jeff Hammerbacher, Information Platforms and the Rise of the Data Scientist. In Beautiful Data, O'Reilly, 2009.

“On the first day of logging the Facebook clickstream, more than 400 gigabytes of data was collected. The load, index, and aggregation processes for this data set really taxed the Oracle data warehouse. Even after significant tuning, we were unable to aggregate a day of clickstream data in less than 24 hours.”

Story for another day...

YAHOO!<sup>®</sup>



# Pig: Example

Find the top 10 most visited pages in each category

Visits

User	Url	Time
Amy	cnn.com	8:00
Amy	bbc.com	10:00
Amy	flickr.com	10:05
Fred	cnn.com	12:00

URL Info

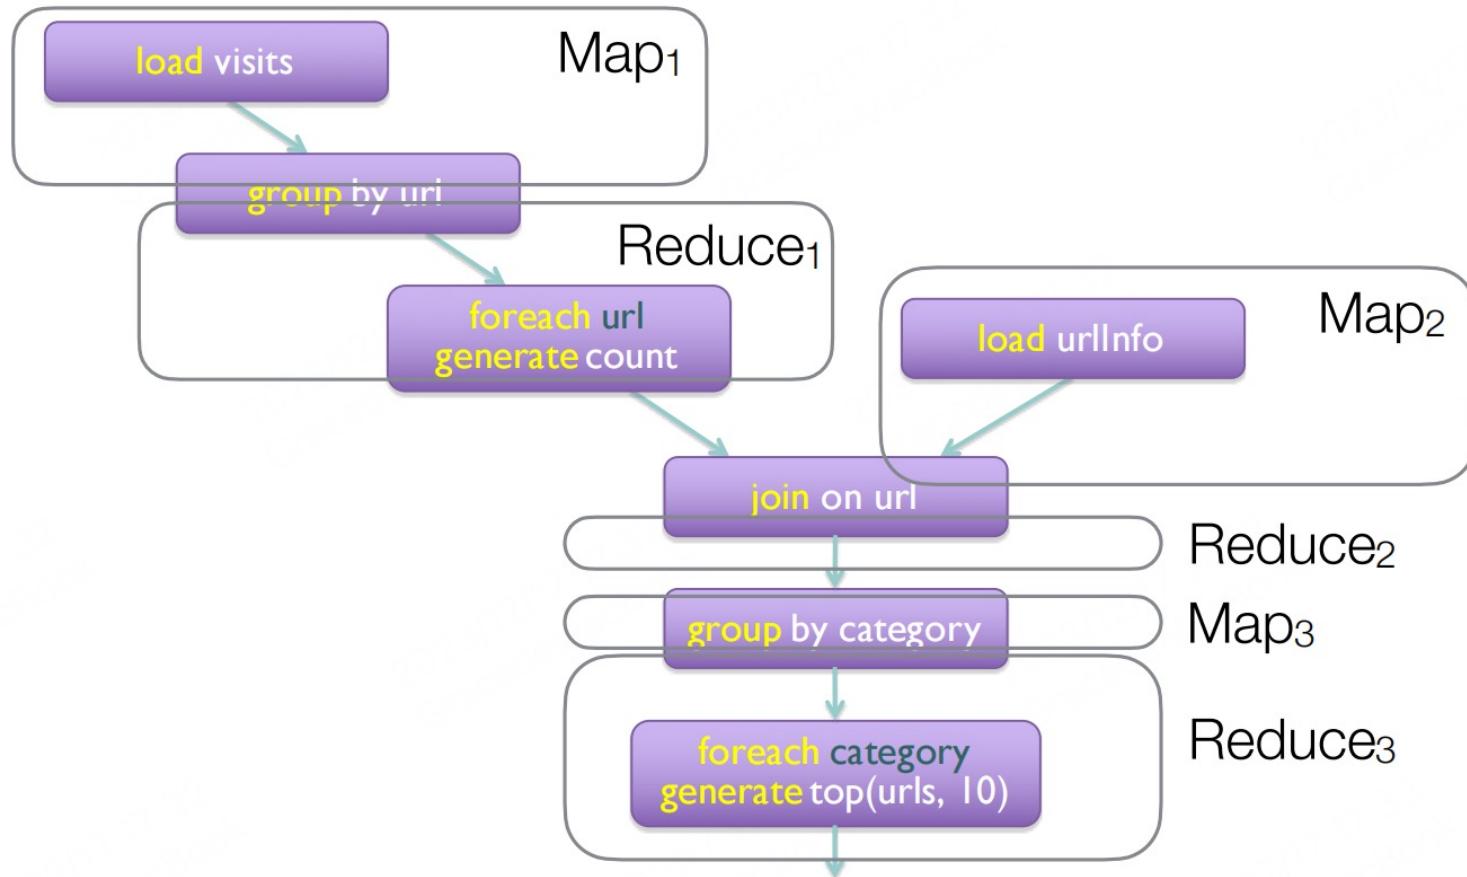
User	Category	PageRank
cnn.com	News	0.9
bbc.com	News	0.8
flickr.com	Photos	0.7
espn.com	Sports	0.9

# Pig: Example

```
visits = load '/data/visits' as (user, url, time);
gVisits = group visits by url;
visitCounts = foreach gVisits generate url, count(visits);
urlInfo = load '/data/urlInfo' as (url, category, pRank);
visitCounts = join visitCounts by url, urlInfo by url;
gCategories = group visitCounts by category;
topUrls = foreach gCategories generate top(visitCounts,10);

store topUrls into '/data/topUrls';
```

# Pig query plan



```

visits = load '/data/visits' as (user, url, time);
gVisits = group visits by url;
visitCounts = foreach gVisits generate url, count(visits);
urlInfo = load '/data/urlInfo' as (url, category, pRank);
visitCounts = join visitCounts by url, urlInfo by url;
gCategories = group visitCounts by category;
topUrls = foreach gCategories generate top(visitCounts,10);

```

Or this?

```

import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.comparator;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.ReduceTask;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.mapred.SequenceFileInputFormat;
import org.apache.hadoop.mapred.SequenceFileOutputFormat;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.jobcontrol.Job;
import org.apache.hadoop.mapred.jobcontrol.JobC ontrol;
import org.apache.hadoop.mapred.lib.IdentityMapper;

public class MRExample {
    public static class LoadPages extends MapReduceBase
        implements Mapper<LongWritable, Text, Text> {
            public void map(LongWritable k, Text val,
                           OutputCollector<Text, Text> oc,
                           Reporter reporter) throws IOException {
                // Pull the key out
                String line = val.toString();
                int firstComma = line.indexOf(',');
                String value = line.substring(firstComma + 1);
                Text outKey = new Text(key);
                // Fprepend an index to the value so we know which file
                // it came from
                Text outVal = new Text("0" + value);
                oc.collect(outKey, outVal);
            }
        }
    public static class LoadAndFilterUsers extends MapReduceBase
        implements Mapper<LongWritable, Text, Text> {
            public void map(LongWritable k, Text val,
                           OutputCollector<Text, Text> oc,
                           Reporter reporter) throws IOException {
                // Pull the key out
                String line = val.toString();
                int firstComma = line.indexOf(',');
                String value = line.substring(firstComma + 1);
                int age = Integer.parseInt(value);
                if (age > 18) {
                    String key = line.substring(0, firstComma);
                    Text outKey = new Text(key);
                    // Fprepend an index to the value so we know which file
                    // it came from
                    Text outVal = new Text("2" + value);
                    oc.collect(outKey, outVal);
                }
            }
        }
    public static class Join extends MapReduceBase
        implements Reducer<Text, Text, Text> {
            public void reduce(Text key,
                              Iterator<Text> iter,
                              OutputCollector<Text, Text> oc,
                              Reporter reporter) throws IOException {
                // For each value, figure out which file it's from and
                store it // accordingly.
                List<String> first = new ArrayList<String>();
                List<String> second = new ArrayList<String>();
                while (iter.hasNext()) {
                    Text t = iter.next();
                    String value = t.toString();
                    if (value.charAt(0) == '1')
                        first.add(value.substring(1));
                    else second.add(value.substring(1));
                }
            }
        }
}

```

This?

```

    reporter.setStatus("OK");
)
// Do the cross product and collect the values
for (String s1 : first) {
    for (String s2 : second) {
        String outval = key + "," + s1 + "," + s2;
        oc.collect(null, new Text(outval));
        reporter.setStatus("OK");
    }
}
}

public static class LoadJoined extends MapReduceBase
    implements Mapper<Text, Text, LongWritable> {
    public void map(
        Text key,
        Text val,
        OutputCollector<Text, LongWritable> oc,
        Reporter reporter) throws IOException {
        // Find the first comma
        String line = val.toString();
        int firstComma = line.indexOf(',');
        String value = line.substring(firstComma + 1);
        // drop the rest of the record, I don't need it anymore,
        // just pass a 1 for the combiner/reducer to sum instead.
        Text outKey = new Text(key);
        oc.collect(outKey, new LongWritable(1));
    }
}

public static class ReduceUrls extends MapReduceBase
    implements Reducer<Text, LongWritable, WritableComparable,
    Writable> {
    public void reduce(
        Text key,
        Iterator<LongWritable> iter,
        OutputCollector<WritableComparable, Writable> oc,
        Reporter reporter) throws IOException {
        long sum = 0;
        while (iter.hasNext()) {
            sum += iter.next().get();
        }
        reporter.setStatus("OK");
    }
    oc.collect(key, new LongWritable(sum));
}

public static class LoadClicks extends MapReduceBase
    implements Mapper<WritableComparable, Writable, LongWritable>, Text> {
    public void map(
        WritableComparable key,
        Writable val,
        OutputCollector<LongWritable, Text> oc,
        Reporter reporter) throws IOException {
        oc.collect((LongWritable)val, (Text)key);
    }
}

public static class LimitClicks extends MapReduceBase
    implements Reducer<LongWritable, Text, LongWritable, Text> {
    int count = 0;
    public void reduce(
        LongWritable key,
        Iterator<Text> iter,
        OutputCollector<LongWritable, Text> oc,
        Reporter reporter) throws IOException {
        // Only output the first 100 records
        while (count < 100 && iter.hasNext()) {
            oc.collect(key, iter.next());
            count++;
        }
    }
}

public static void main(String[] args) throws IOException {
    JobConf ip = new JobConf(MRExample.class);
    ip.setJobName("Load Pages");
    ip.setInputFormat(TextInputFormat.class);
    Path[] user Gates Pages = new Path[1];
    user Gates Pages[0] = new Path("/user/gates/pages");
    ip.setNumReduceTasks(0);
    Job loadPages = new Job(ip);
    JobConf ltu = new JobConf(MRExample.class);
    ltu.setJobName("Load and Filter Users");
    ltu.setInputFormat(TextInputFormat.class);
    ltu.setOutputFormat(TextOutputFormat.class);
    ltu.setMapperClass(LoadAndFilterUsers.class);
    ltu.setReducerClass(LoadAndFilterUsers.class);
    FileInputFormat.addInputPath(ltu, new Path("/user/gates/filtered"));
    FileOutputFormat.setOutputPath(ltu,
        new Path("/user/gates/tmp/filtered_users"));
    ltu.setNumReduceTasks(0);
    Job loadUsers = new Job(ltu);
    JobConf join = new JobConf(MRExample.class);
    join.setJobName("Join User and Pages");
    join.setInputFormat(TextInputFormat.class);
    join.setOutputValueClass(Text.class);
    join.setOutputKeyClass(Text.class);
    join.setMapperClass(Join.class);
    join.setReducerClass(Join.class);
    FileInputFormat.addInputPath(join, new Path("/user/gates/tmp/indexed_pages"));
    FileInputFormat.addInputPath(join, new Path("/user/gates/tmp/filtered_users"));
    FileOutputFormat.setOutputPath(join, new Path("/user/gates/tmp/joined"));
    join.setNumReduceTasks(1);
    Job joinJob = new Job(join);
    joinJob.addDependingJob(loadPages);
    joinJob.addDependingJob(loadUsers);
    joinJob.setGroup(group);
    joinJob.setJobName("Group URLs");
    joinJob.setMapperClass(KeyValueInputFormat.class);
    joinJob.setReducerClass(SequenceFileOutputFormat.class);
    joinJob.setCombinerClass(ReduceUrls.class);
    joinJob.setPartitionerClass(SequenceFilePartitioner.class);
    FileInputFormat.addInputPath(join, new Path("/user/gates/tmp/joined"));
    FileInputFormat.setOutputPath(join, new Path("/user/gates/tmp/grouped"));
    joinJob.setNumReduceTasks(50);
    Job groupJob = new Job(join);
    groupJob.addDependingJob(joinJob);
    JobConf top100 = new JobConf(MRExample.class);
    top100.setJobName("Top 100 sites");
    top100.setInputFormat(SequenceFileInputFormat.class);
    top100.setMapperClass(LoadClicks.class);
    top100.setReducerClass(LimitClicks.class);
    top100.setPartitionerClass(SequenceFilePartitioner.class);
    top100.setNumReduceTasks(1);
    FileInputFormat.addInputPath(top100, new Path("/user/gates/tmp/grouped"));
    Path[] top100sitesforusers = new Path[1];
    top100sitesforusers[0] = new Path("top100sitesforusers");
    top100.setNumReduceTasks(1);
    Job limit = new Job(top100);
    limit.addDependingJob(groupJob);
    JobControl jc = new JobControl("Find top
100 sites for users
18 to 25");
    jc.addJob(loadPages);
    jc.addJob(loadUsers);
    jc.addJob(joinJob);
    jc.addJob(groupJob);
    jc.addJob(limit);
    jc.run();
}

```

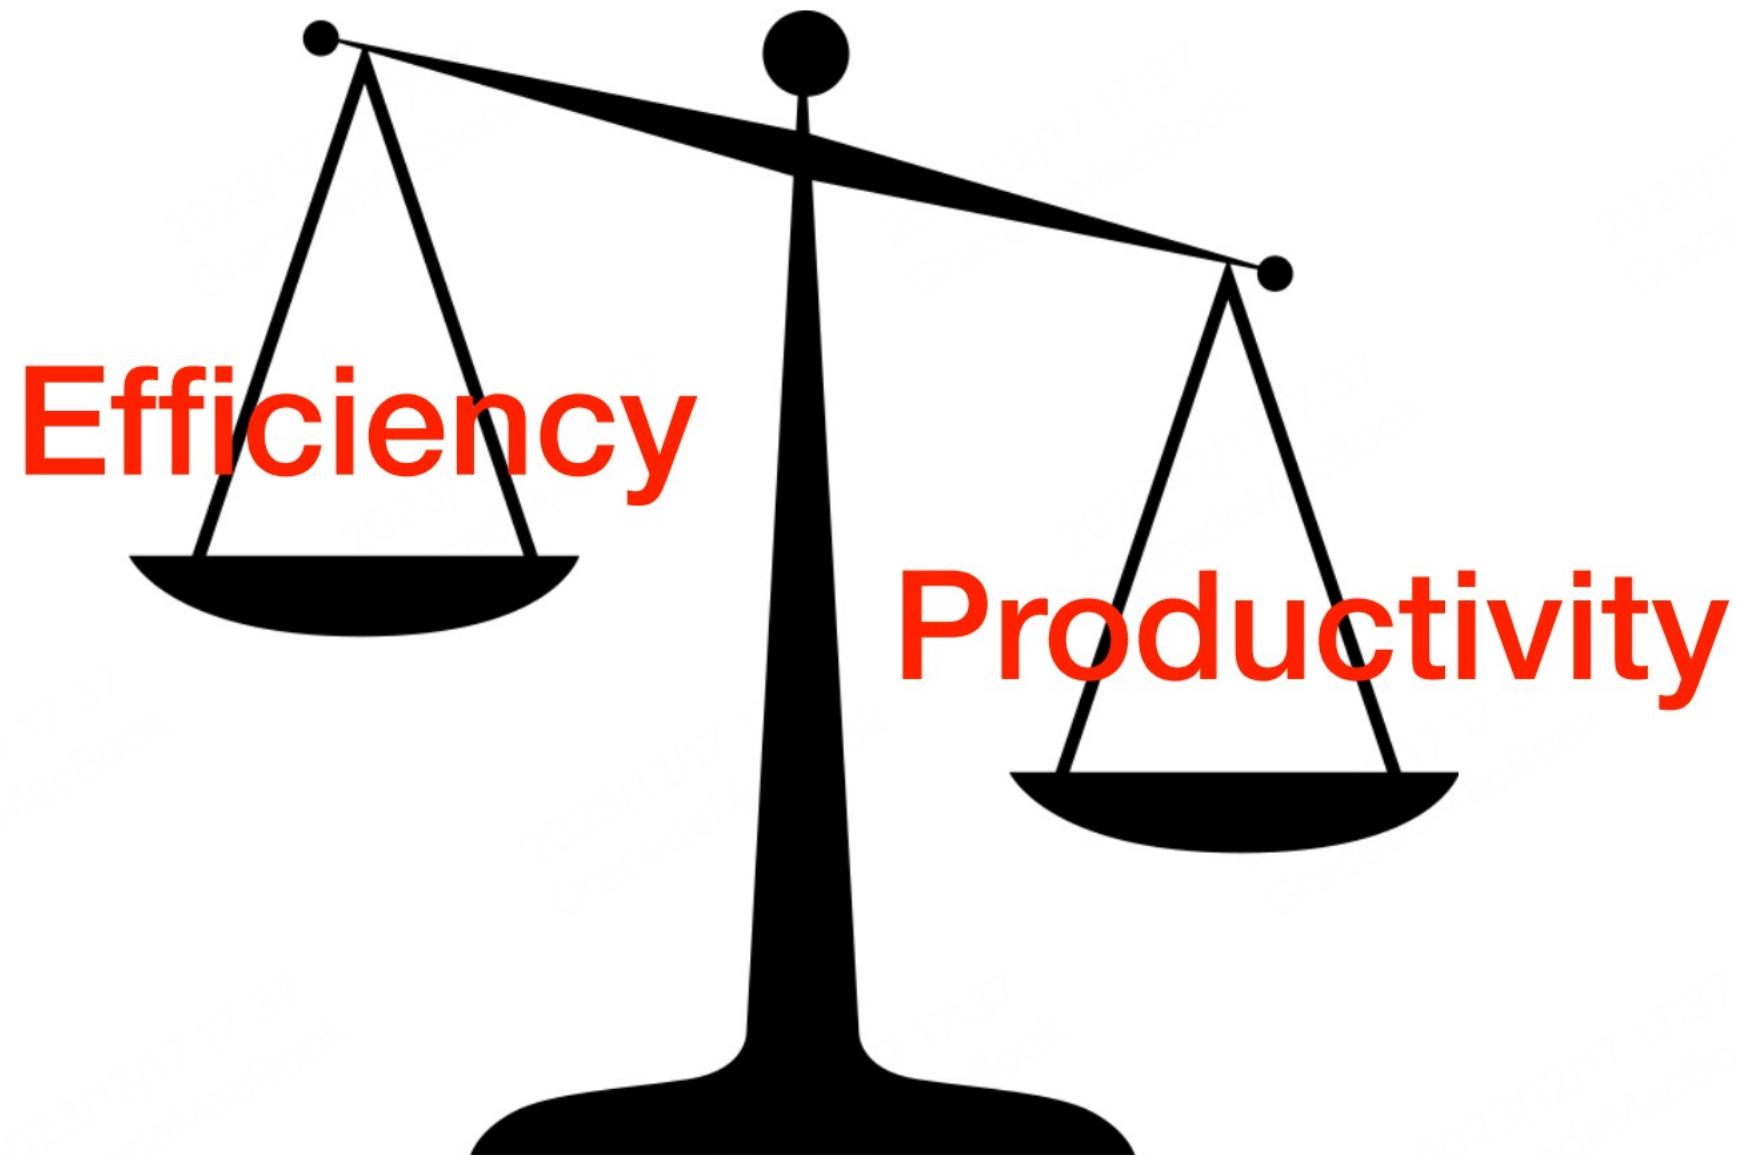
# Programming at a high level

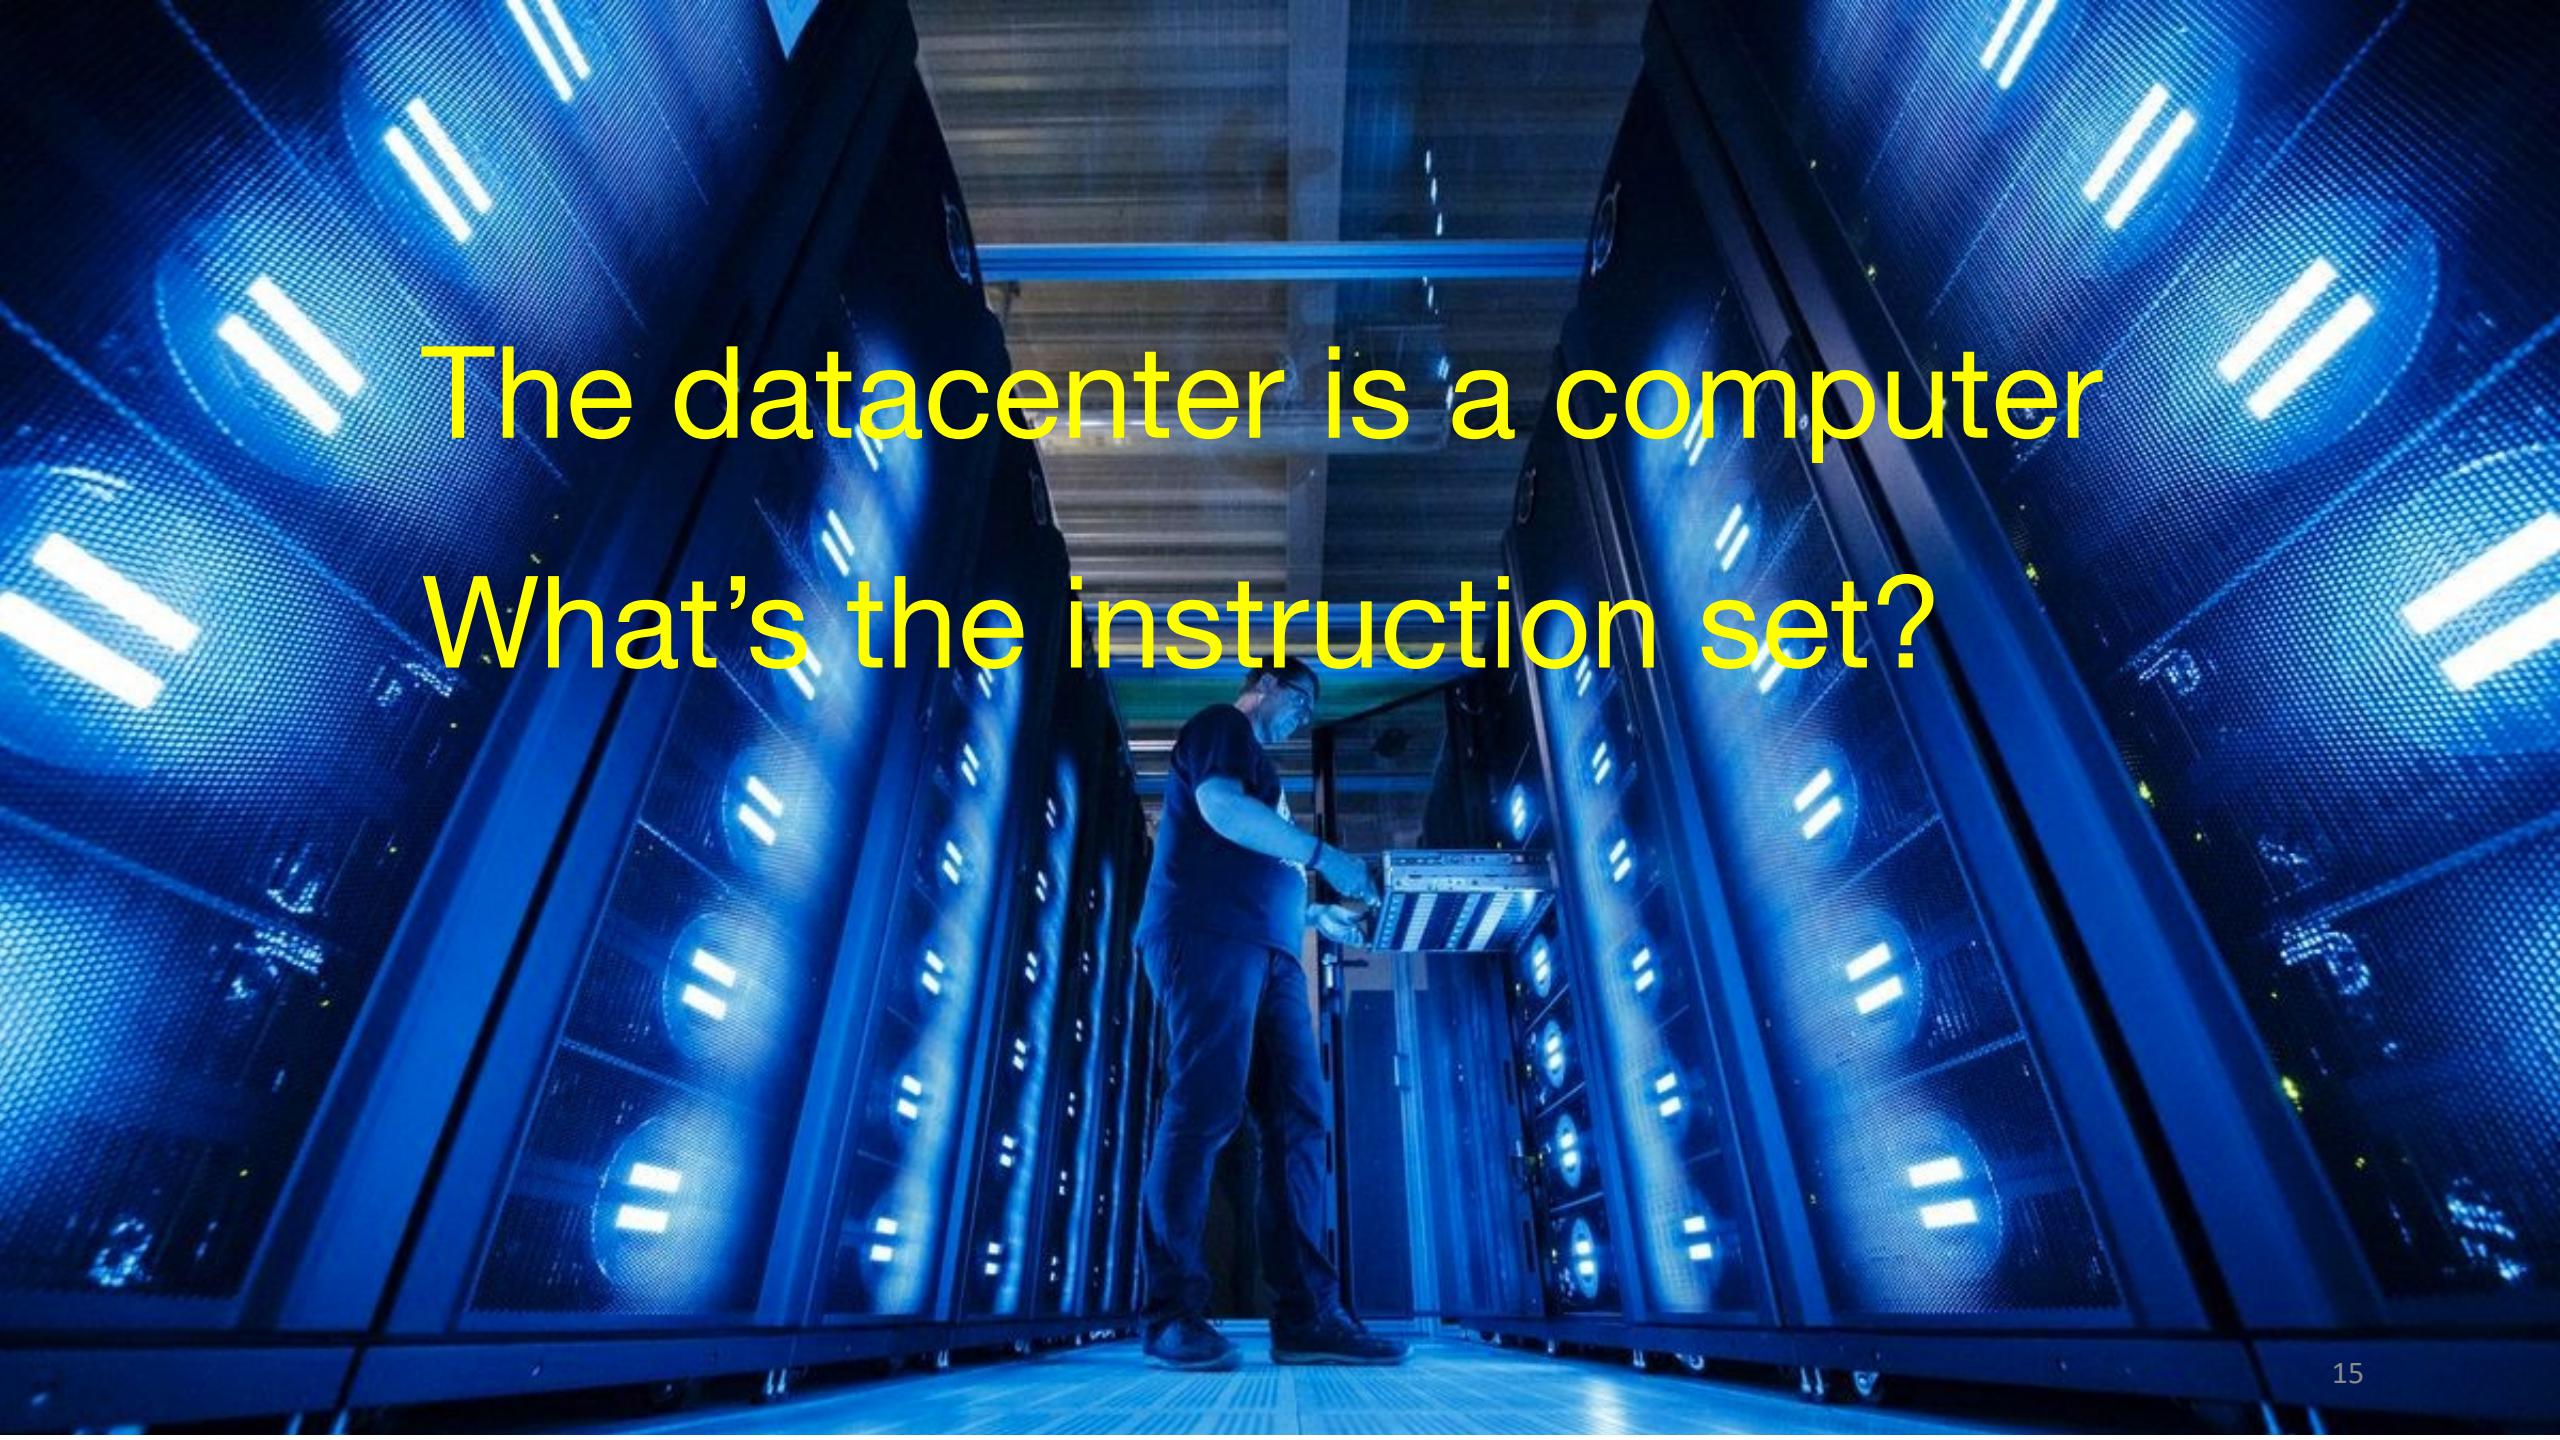
Specify “**what**”, and let the system figure out  
“**how**”

# But isn't Pig slower?

Sure, but C can be slower than assembly too...

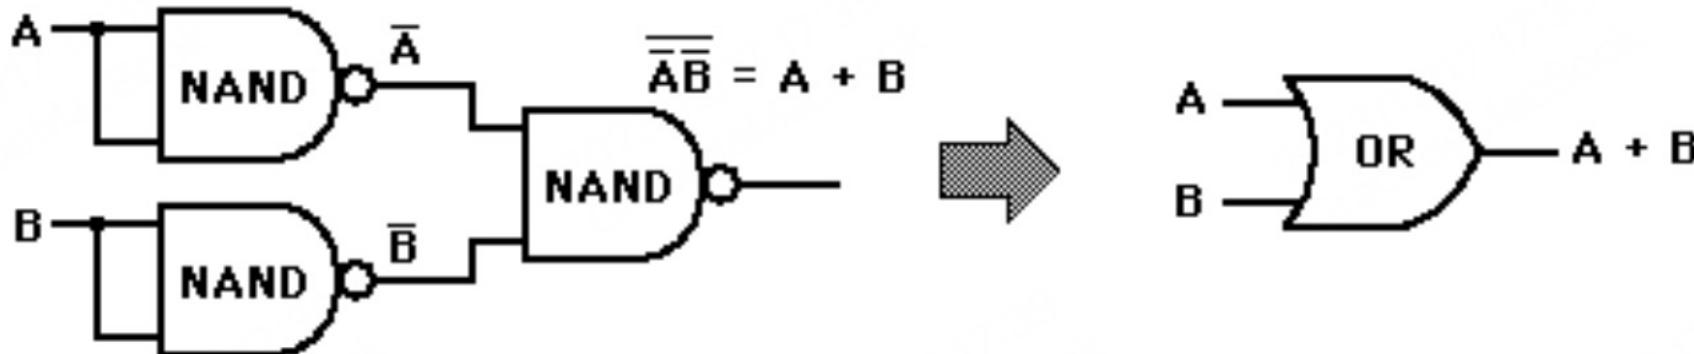
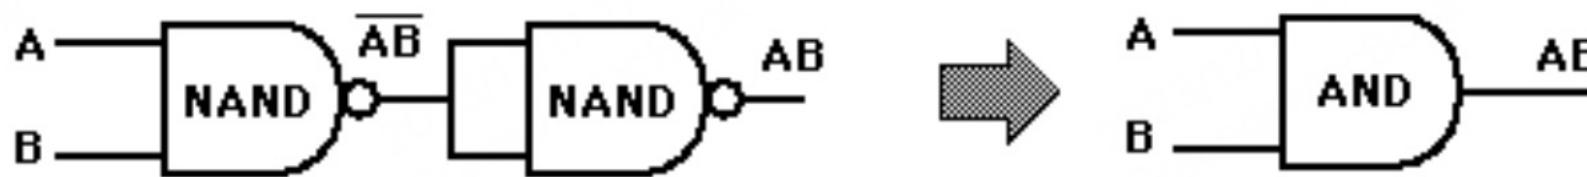
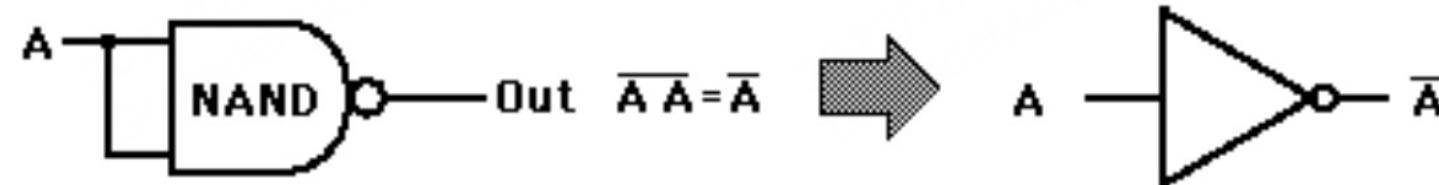




A photograph of a datacenter aisle. A person in a dark t-shirt and pants is standing in the center, facing a row of server racks. The racks are illuminated from within, casting a blue glow. The ceiling has recessed lighting, and the floor is a polished blue-tinted surface.

The datacenter is a computer  
What's the instruction set?

# Analogy: NAND Gates are universal



# Let's design a data processing language “from scratch”!

Why is MapReduce the way it is?

# Data-parallel dataflow languages

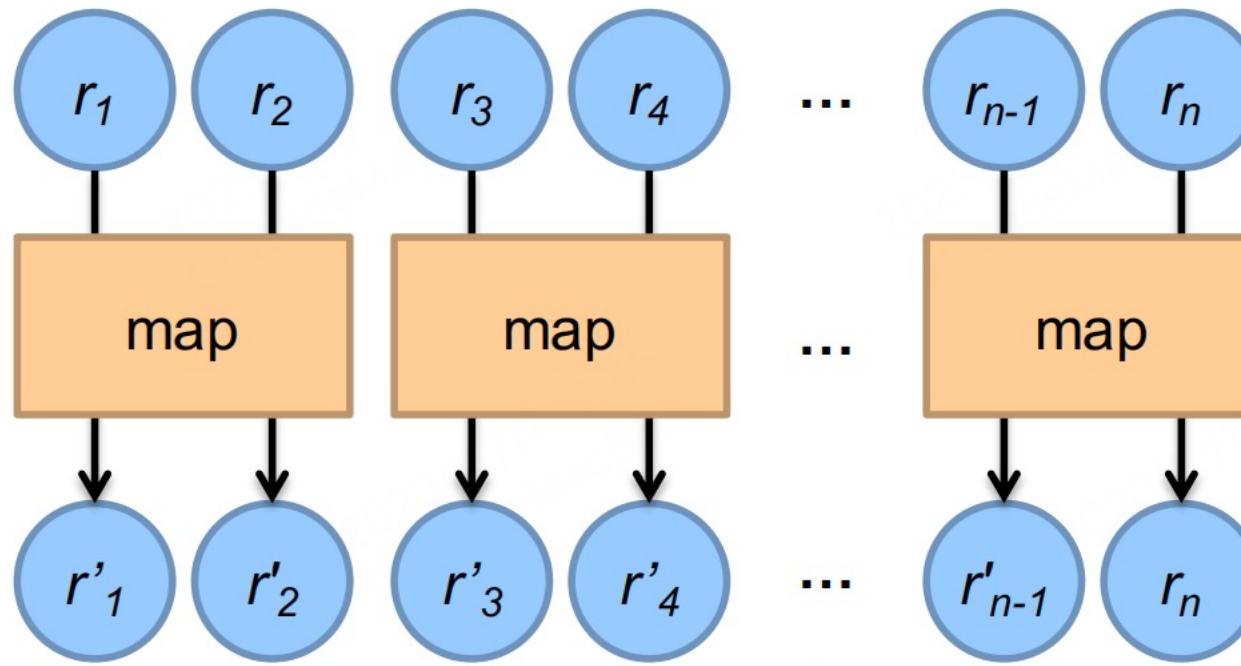
We have a collection of **records**

- want to apply a bunch of transformations to compute some results

Assumptions

- static collection records (not necessarily key-value pairs)

We need per-record processing  
(note: not necessarily key-value pairs)



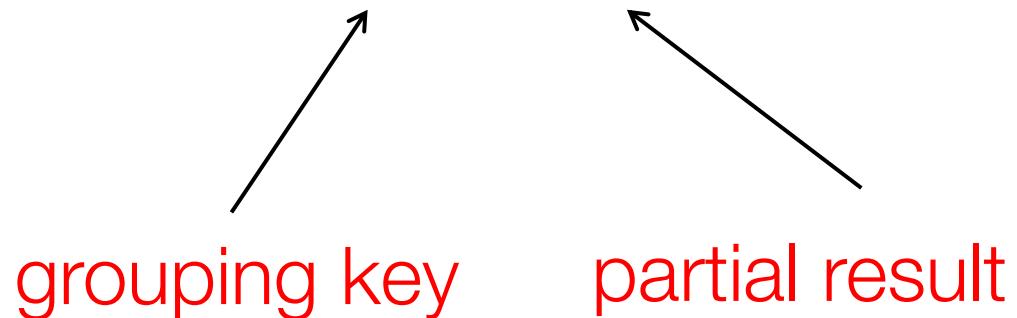
Easy to parallelize maps

Record to “mapper” assignment is an implementation detail

# Map alone isn't enough

We need a way to group partial results

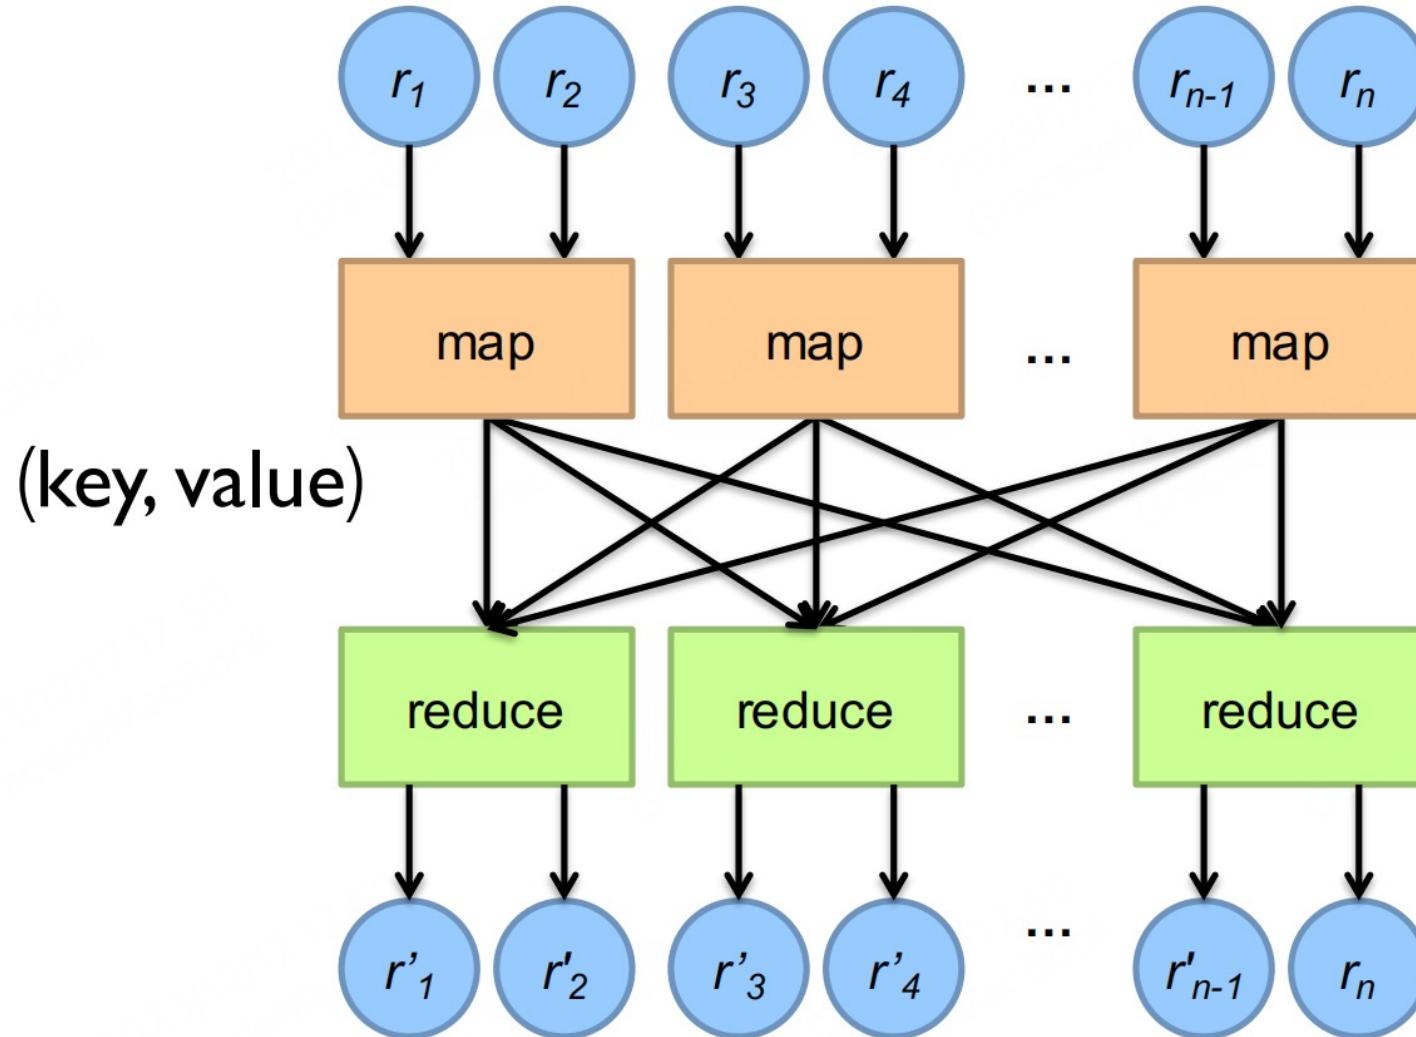
- intermediate **(key, value)** pairs



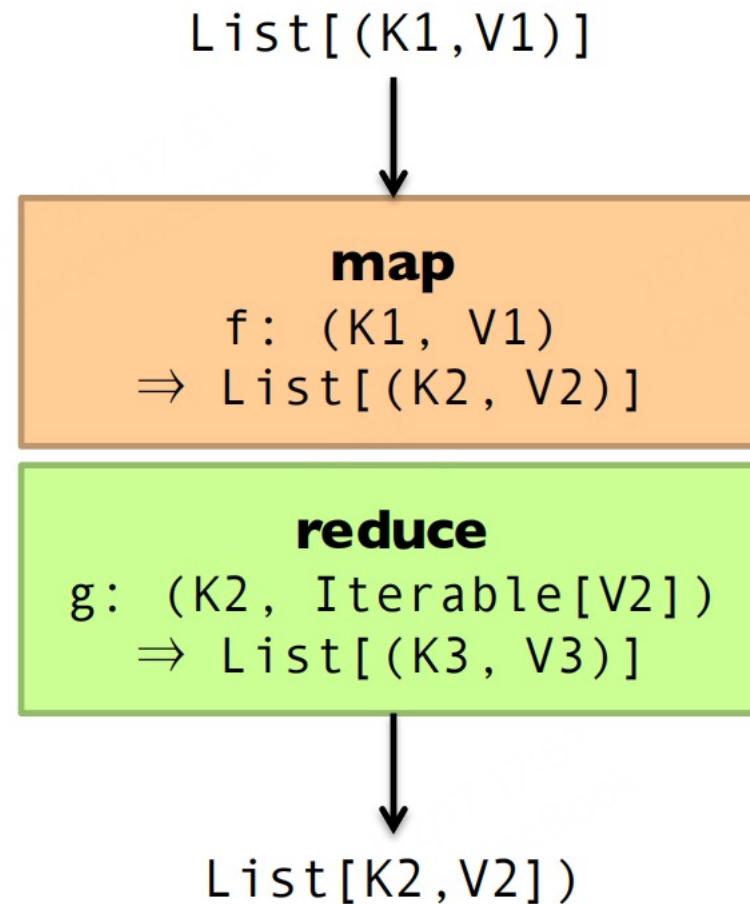
For each key, we can apply some computation

- e.g., aggregation, counting

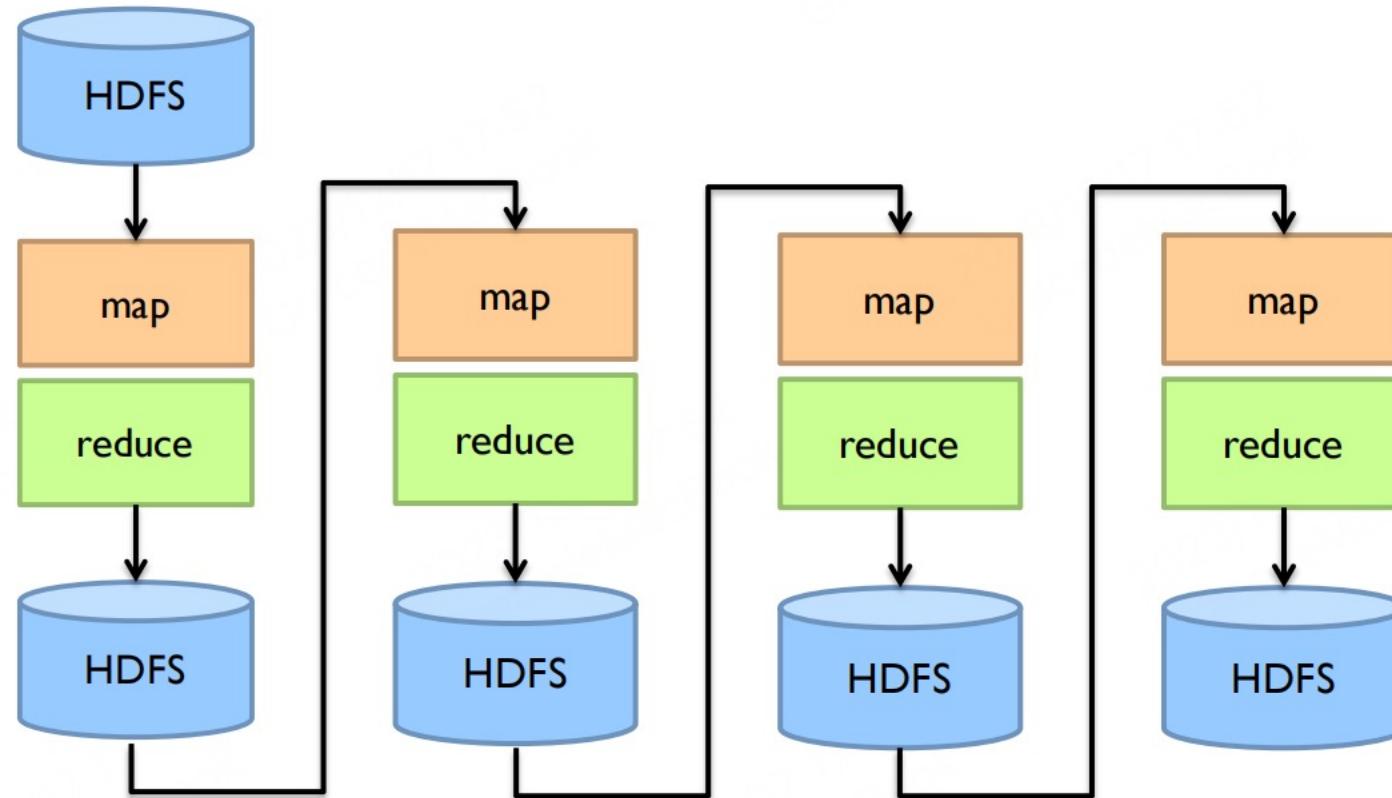
# MapReduce



# MapReduce

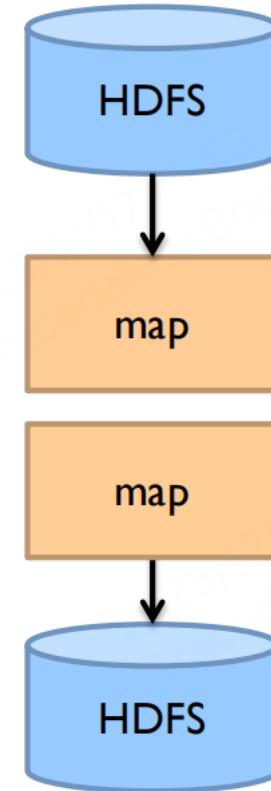
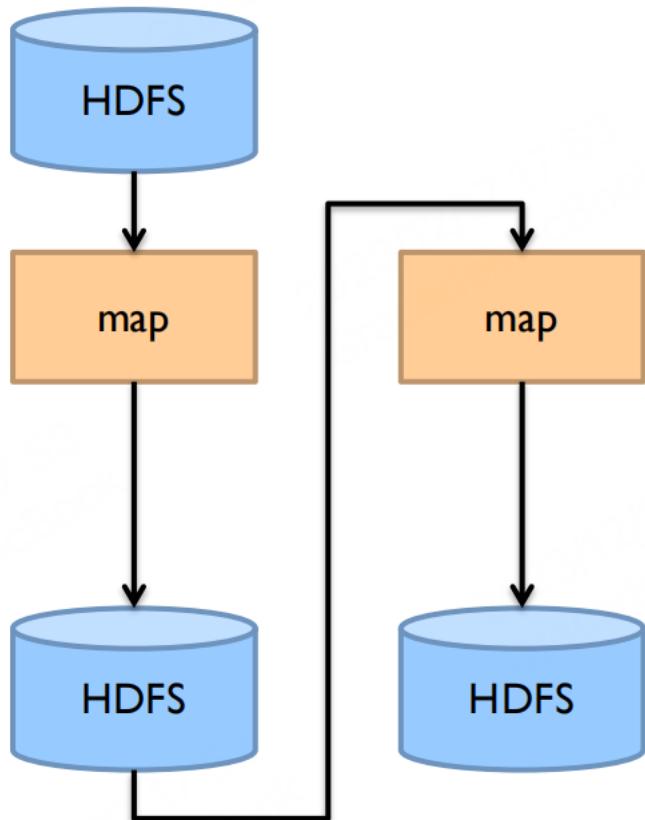


# MapReduce workflows

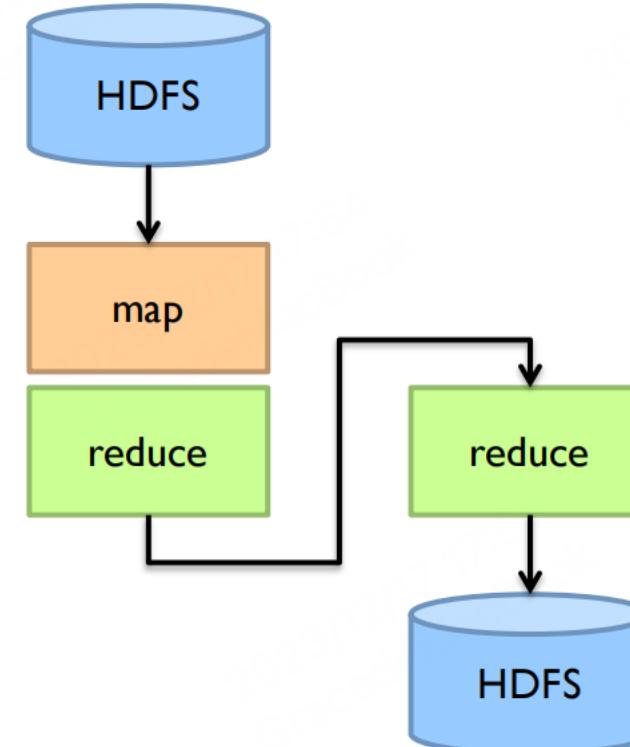
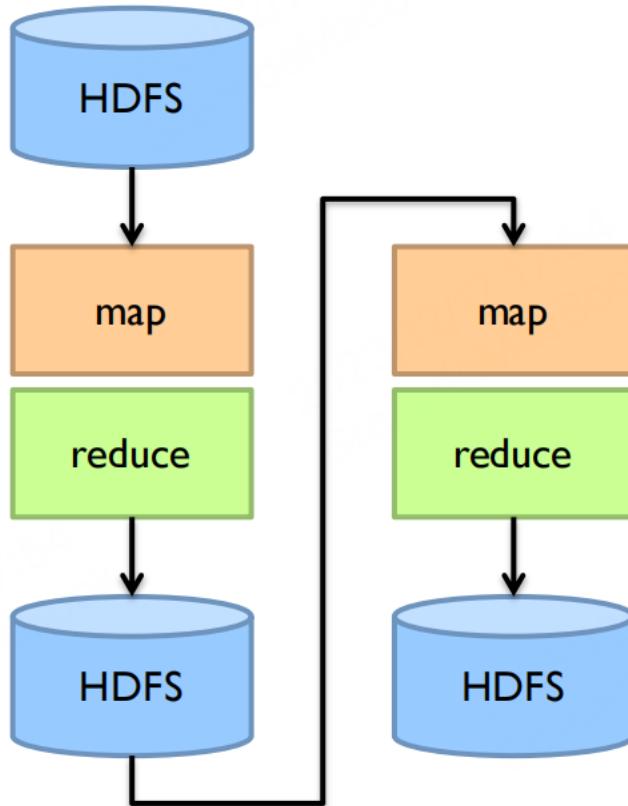


What's wrong?

# Want MM?



# Want MRR?



In MapReduce, the only way to share data across jobs is stable storage

Disk I/O is waaaaaaay too slow!

# Data-parallel dataflow languages

We have a collection of **records**

- want to apply a bunch of transformations to compute some results

What are the operators?

- MapReduce?





Where the hype is!

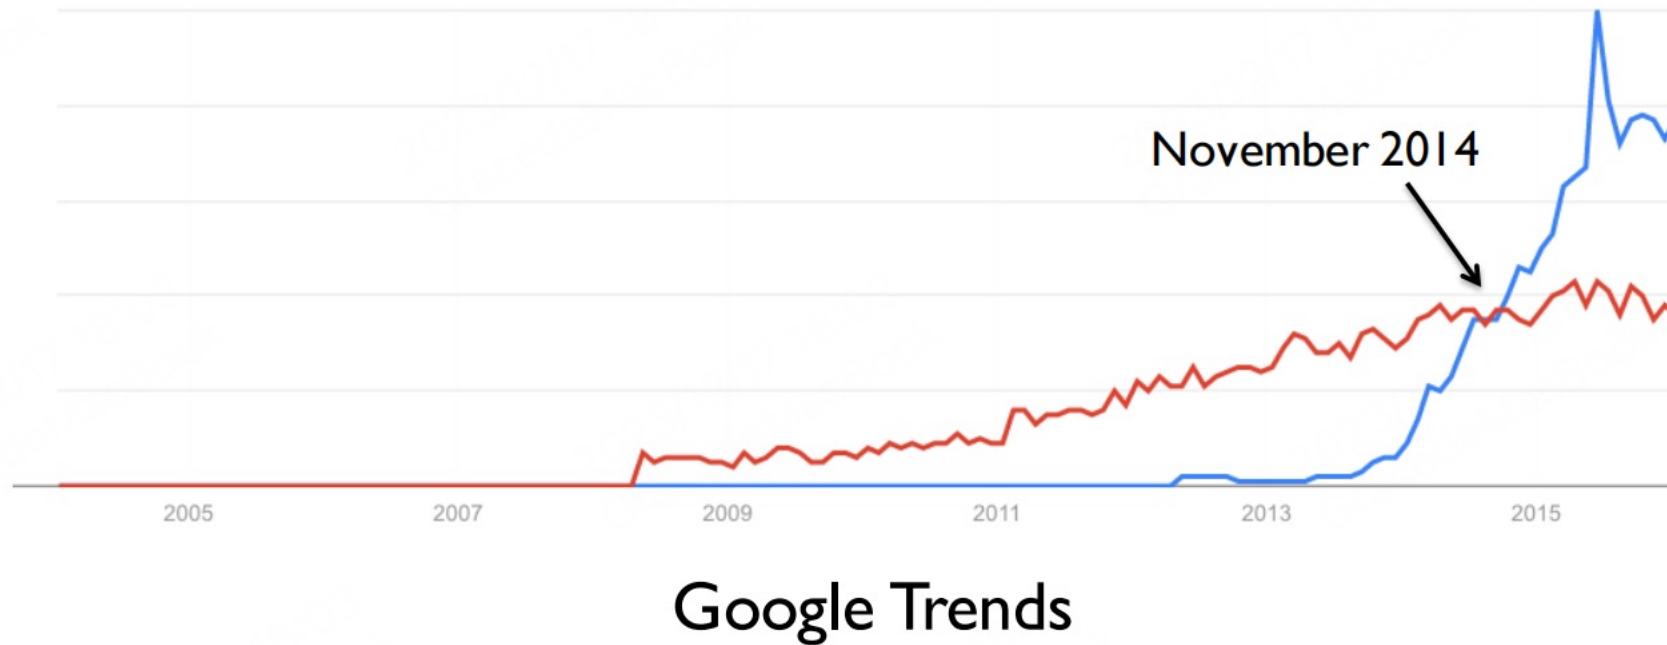
- answer to “What’s beyond MapReduce?”

Brief history:

- developed at UC Berkeley AMPLab in 2009
- open-sourced in 2010
- became top-level Apache project in Feb 2014
- commercial support provided by Databricks



vs.



Source: Datanami (2014): <http://www.datanami.com/2014/11/21/spark-just-passed-hadoop-popularity-web-heres/>

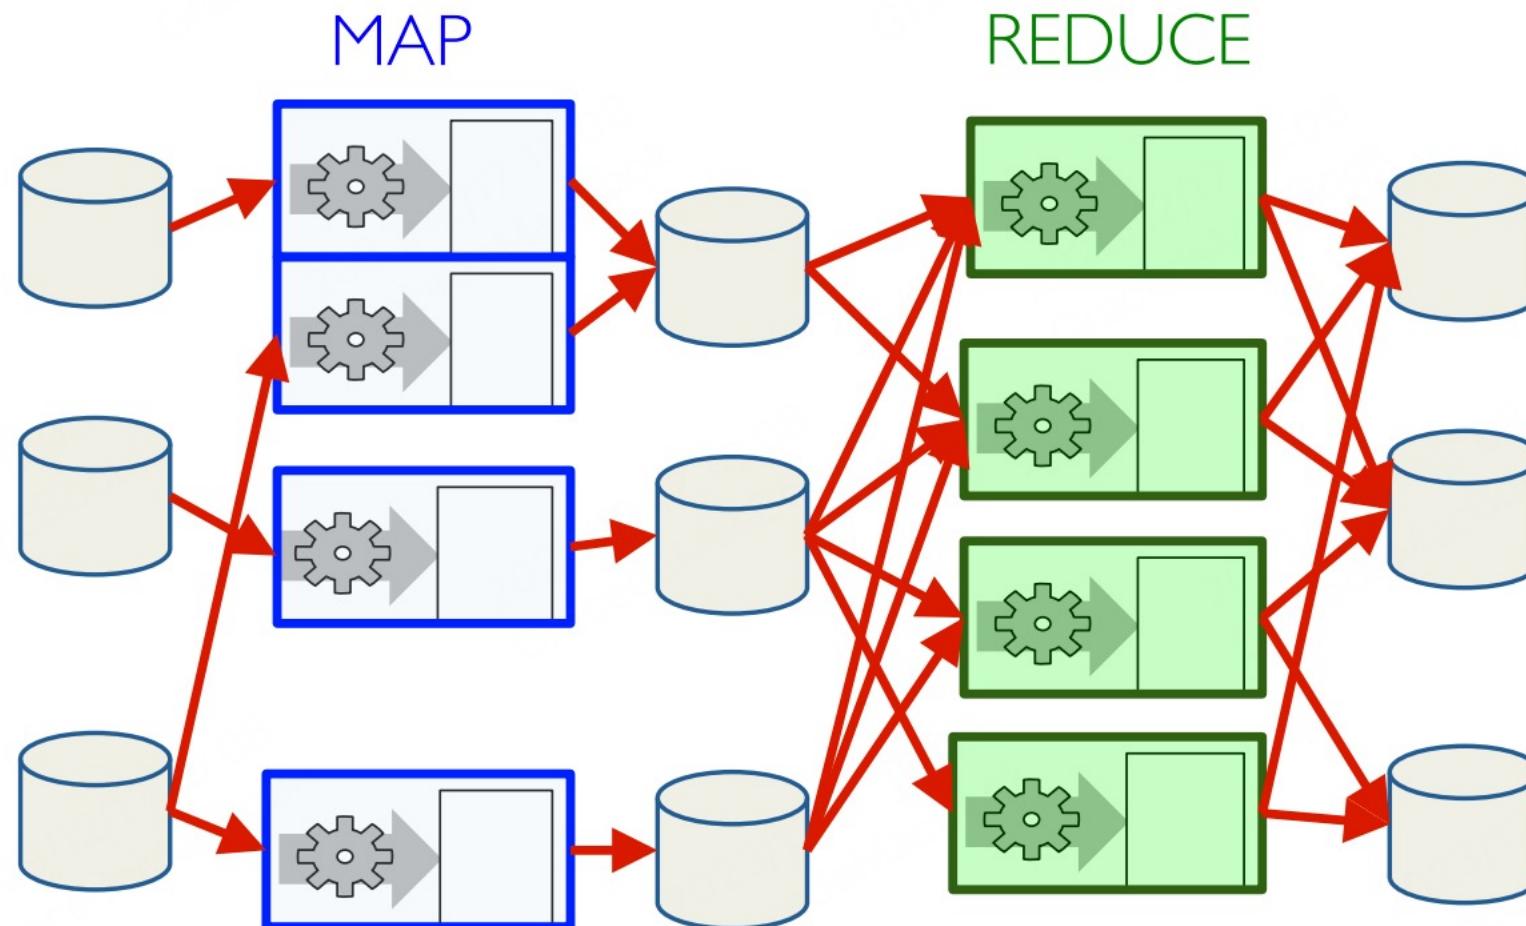
# Motivation

MapReduce is a restrictive programming model for batch processing jobs

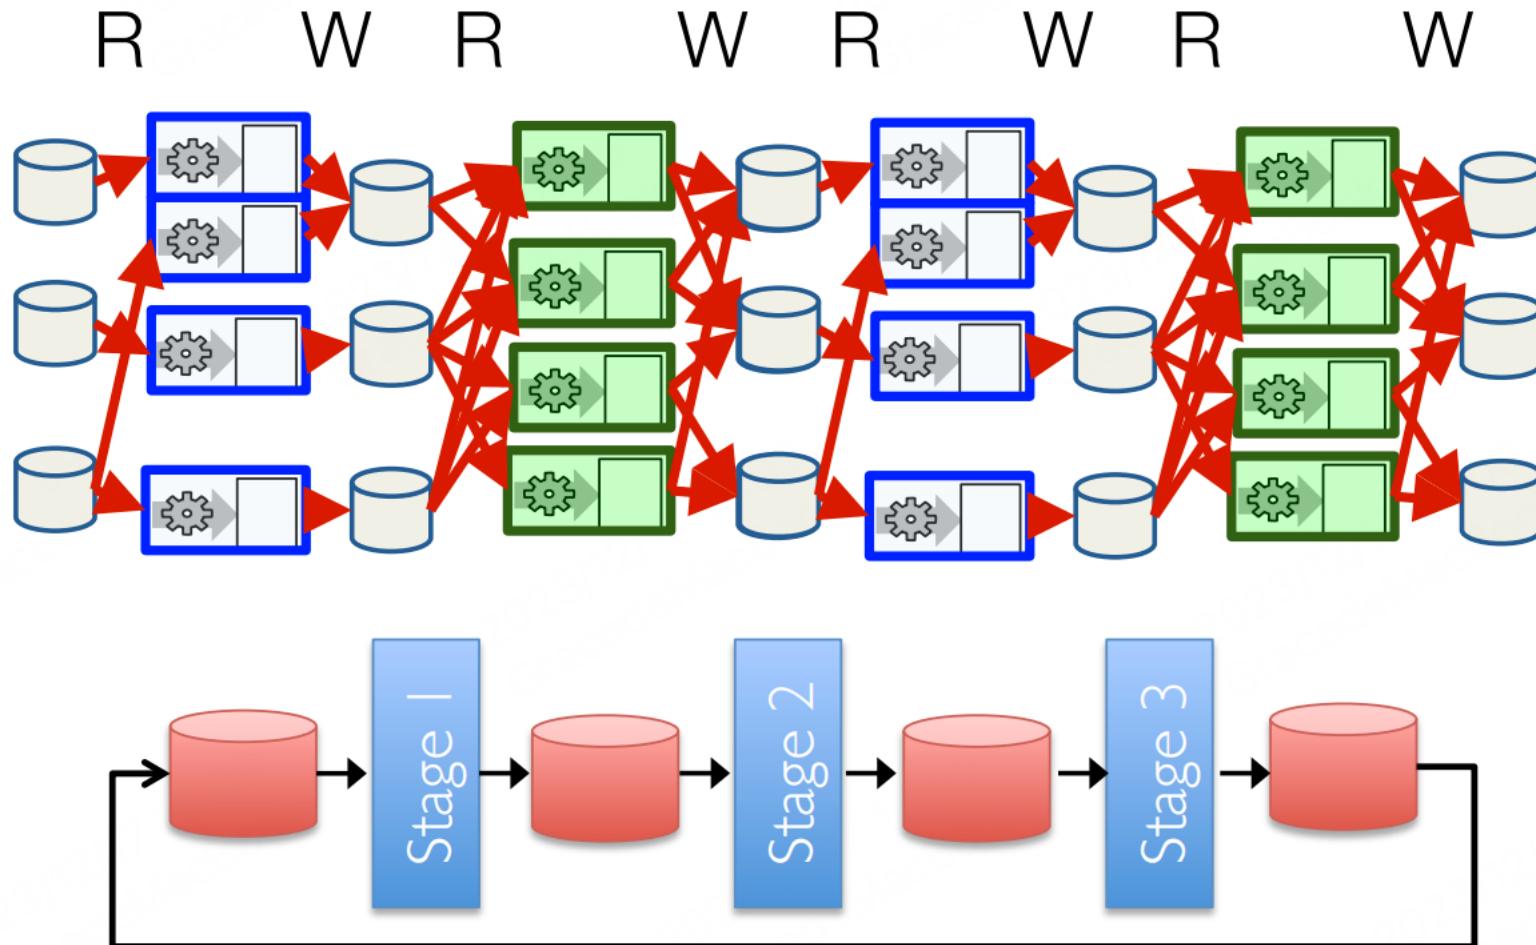
But we want more...

- More **complex**, multi-stage data processing, e.g., iterative machine learning
- Low latency, **interactive** ad-hoc data analytics, e.g., streaming processing

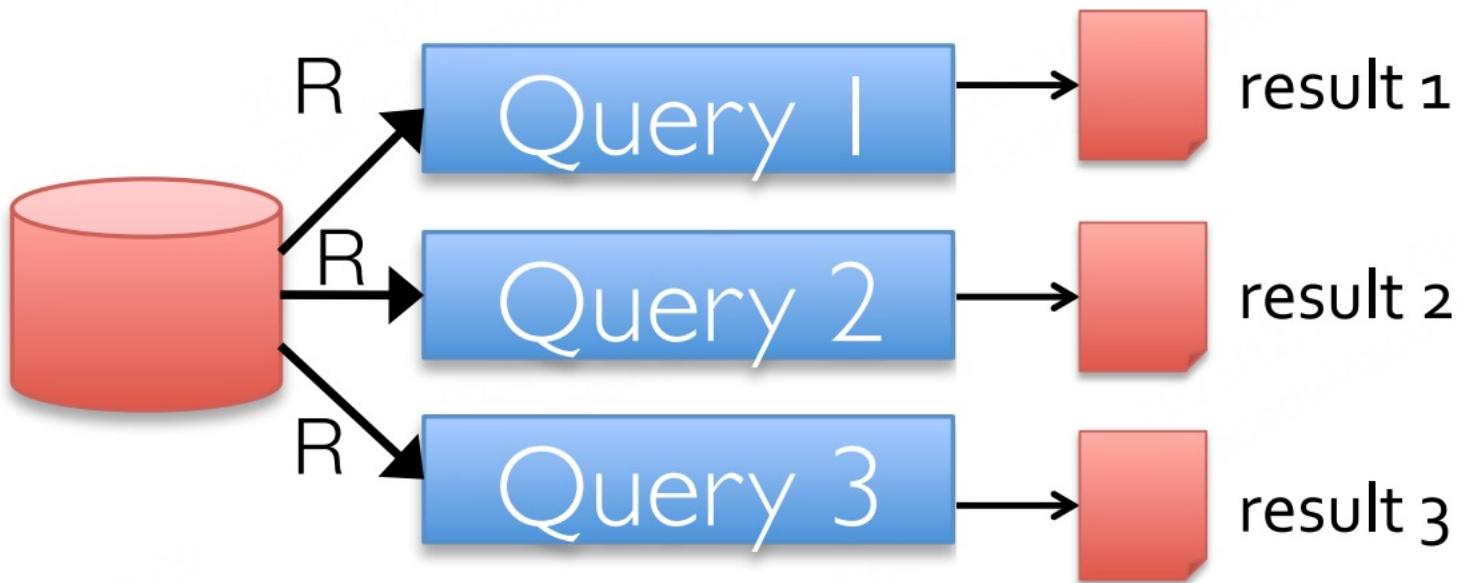
Each stage passes through the stable storage!



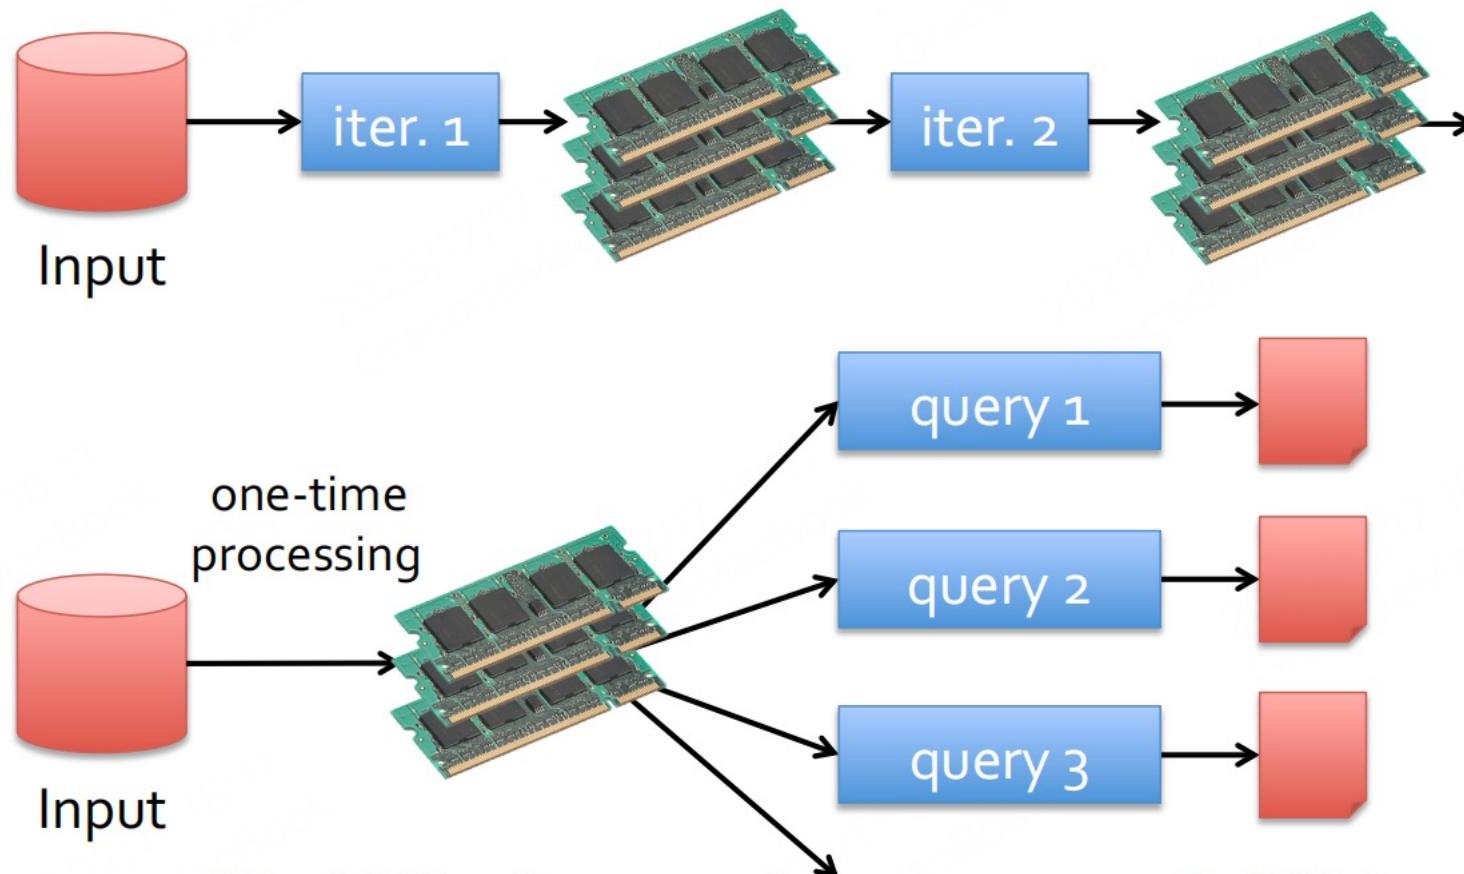
# Iterative computations



# Iterative queries

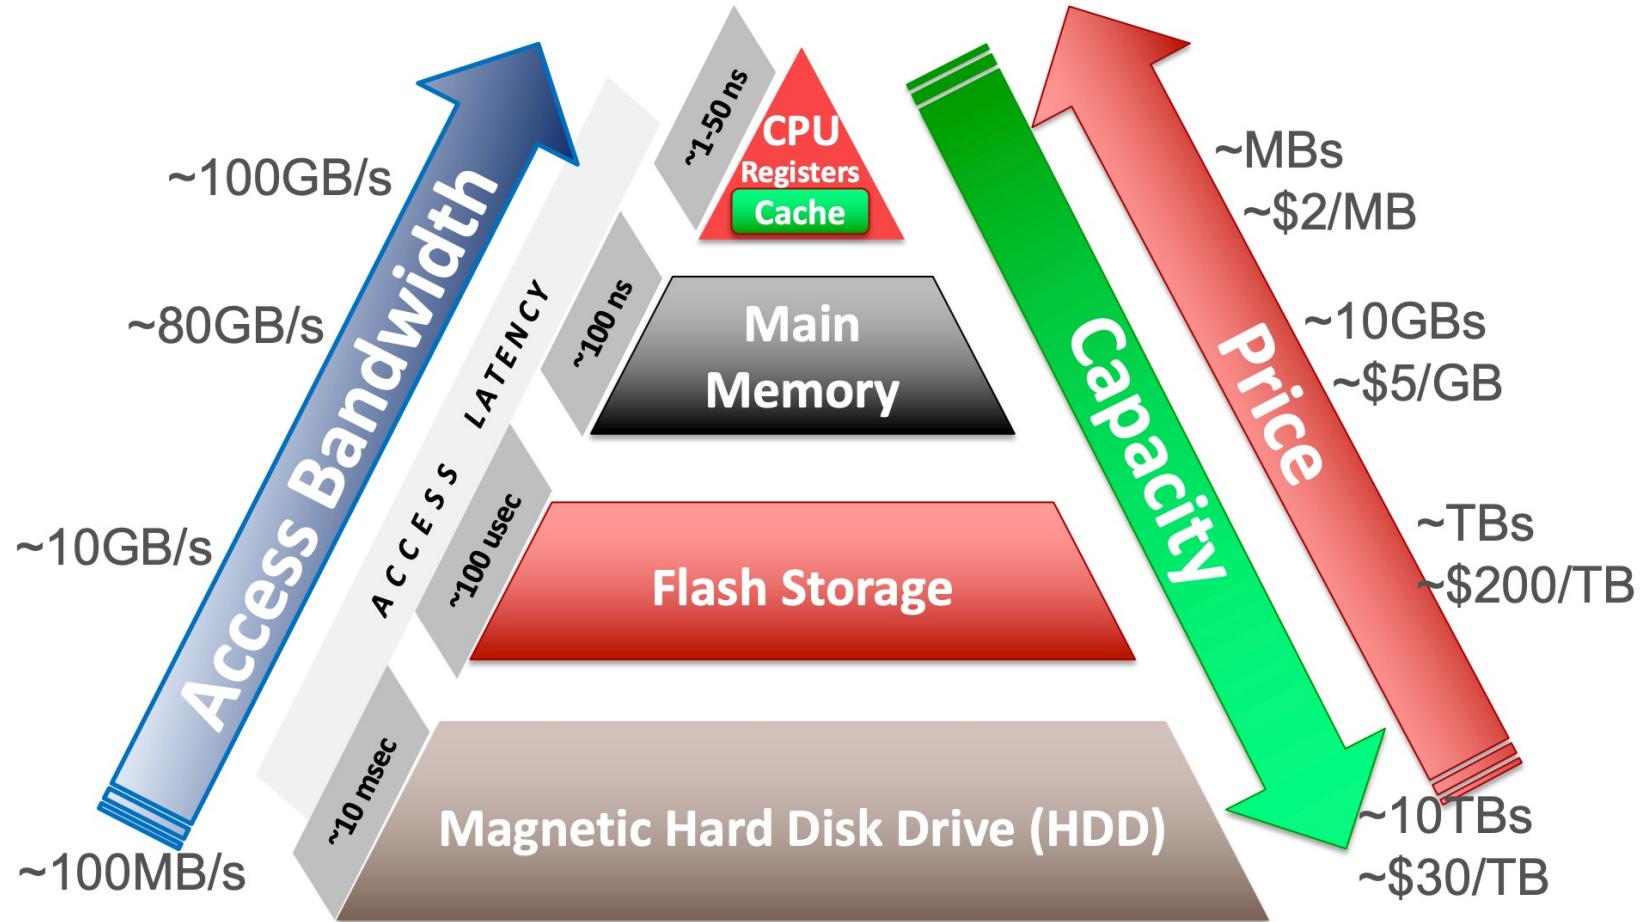


# Intuition: In-memory data sharing

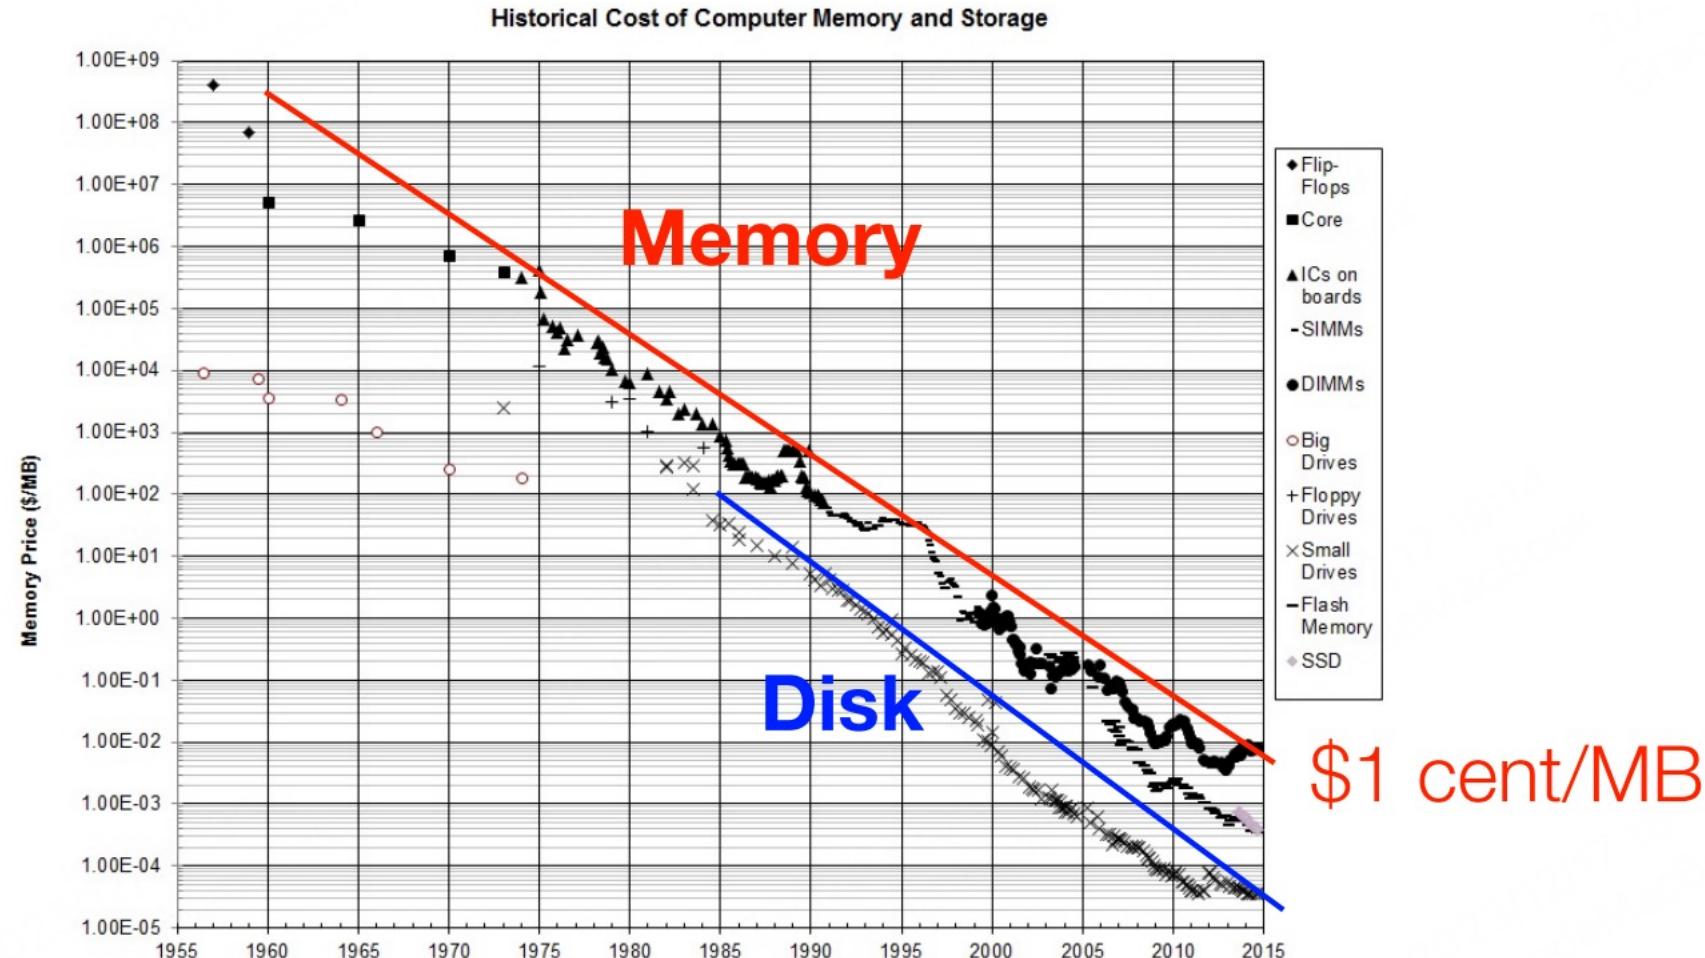


10-100x faster than network/disk

# Memory-storage hierarchy



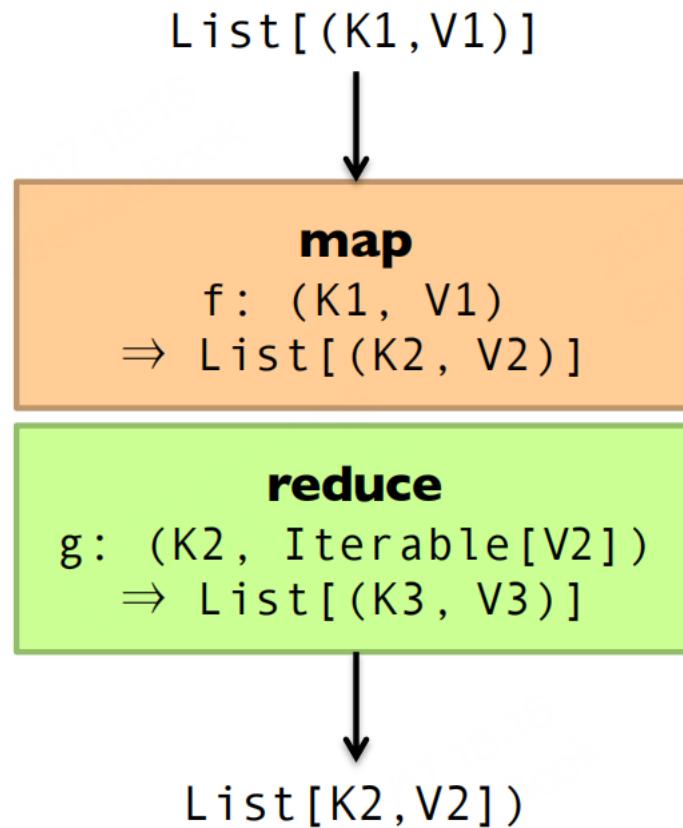
# Cost of Memory



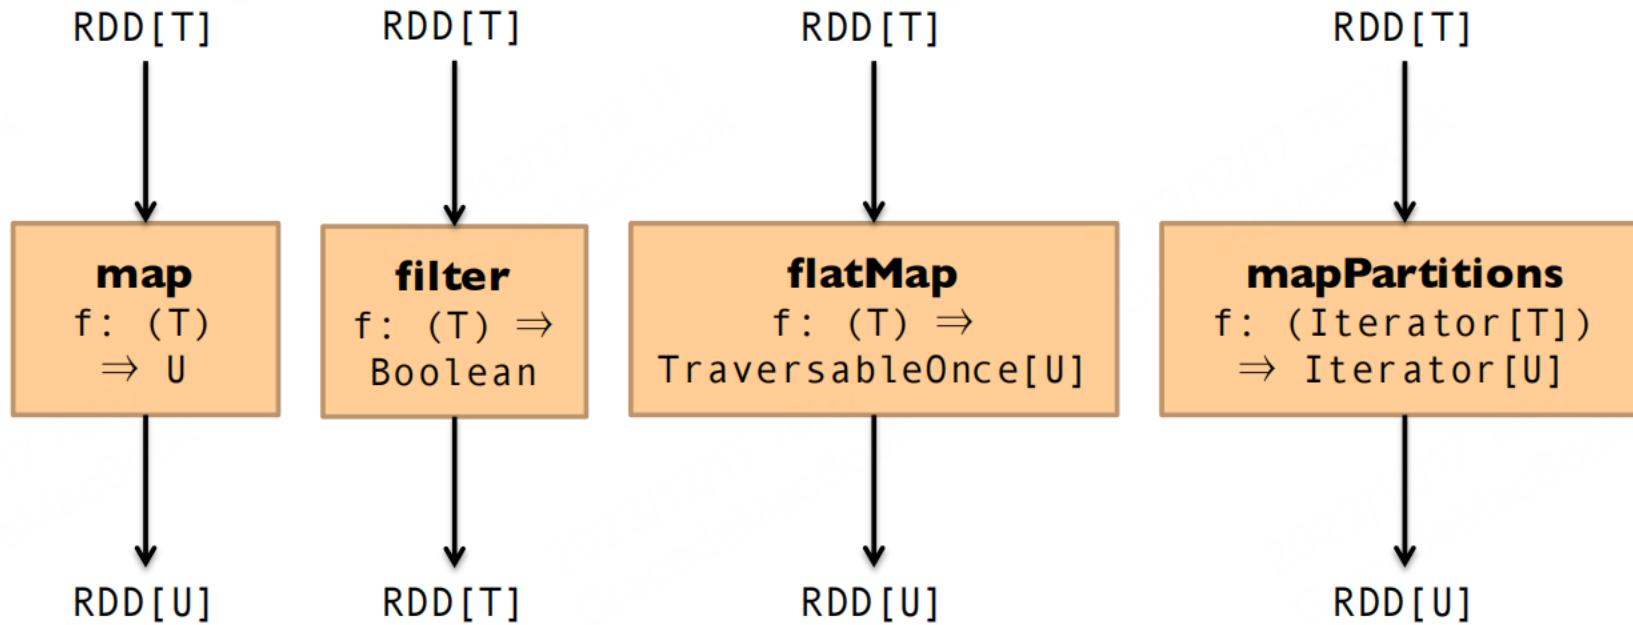
Intermediate results are stored in-memory  
as **Resilient Distributed Datasets** (RDDs)

Much more next lecture...

# MapReduce

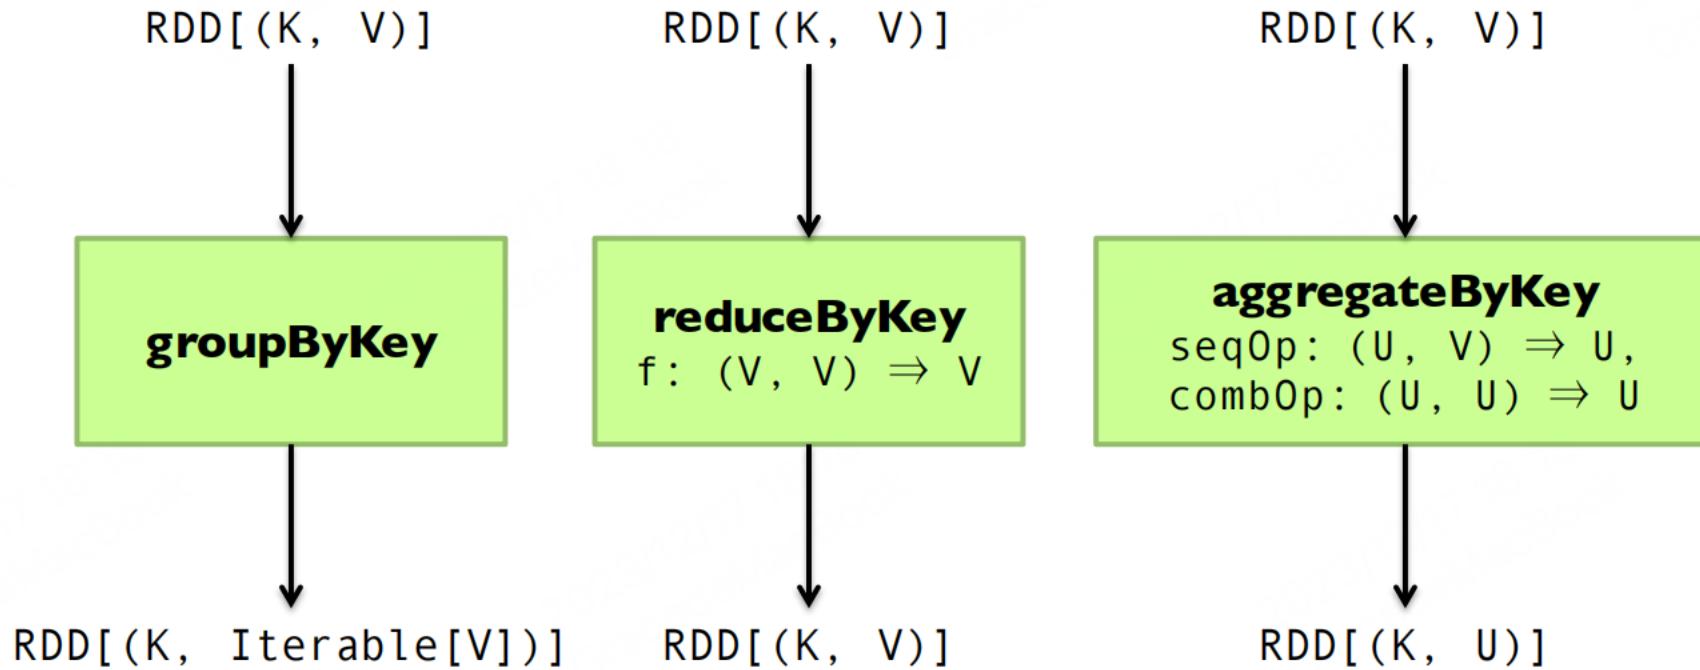


# Map-like operations



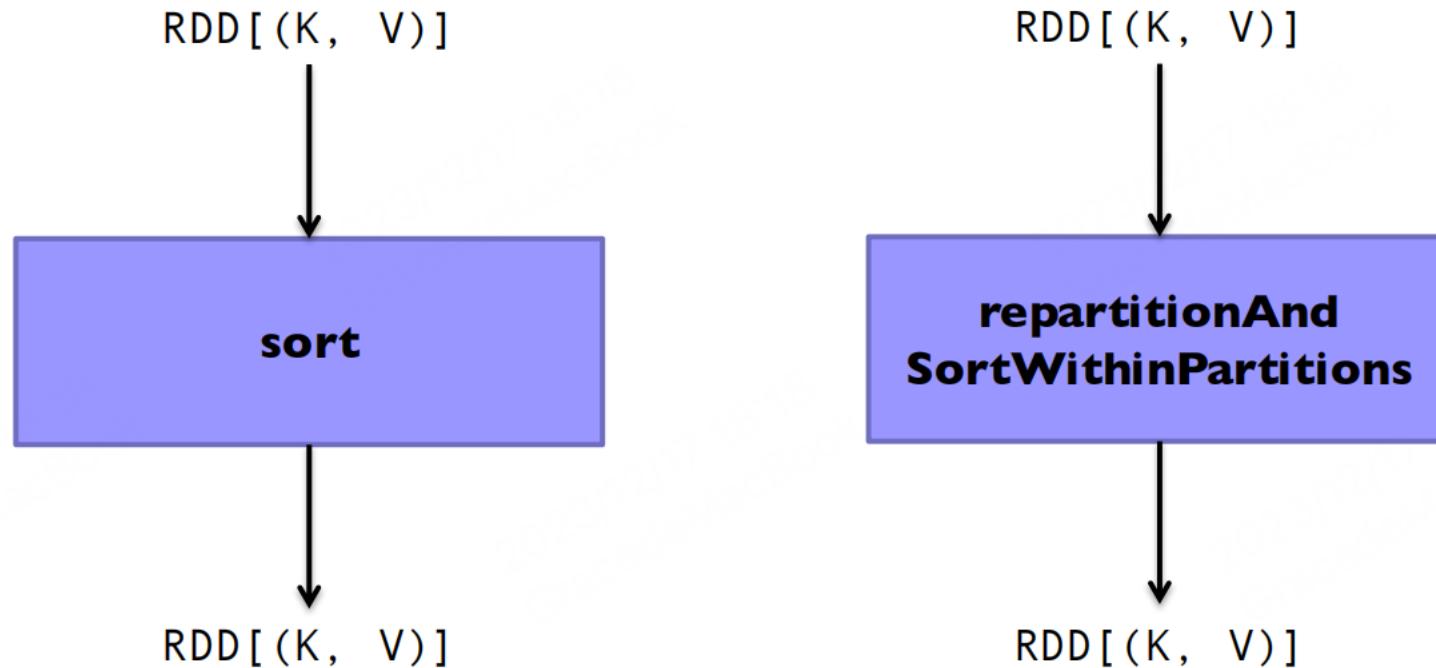
Not meant to be exhaustive

# Reduce-like operations



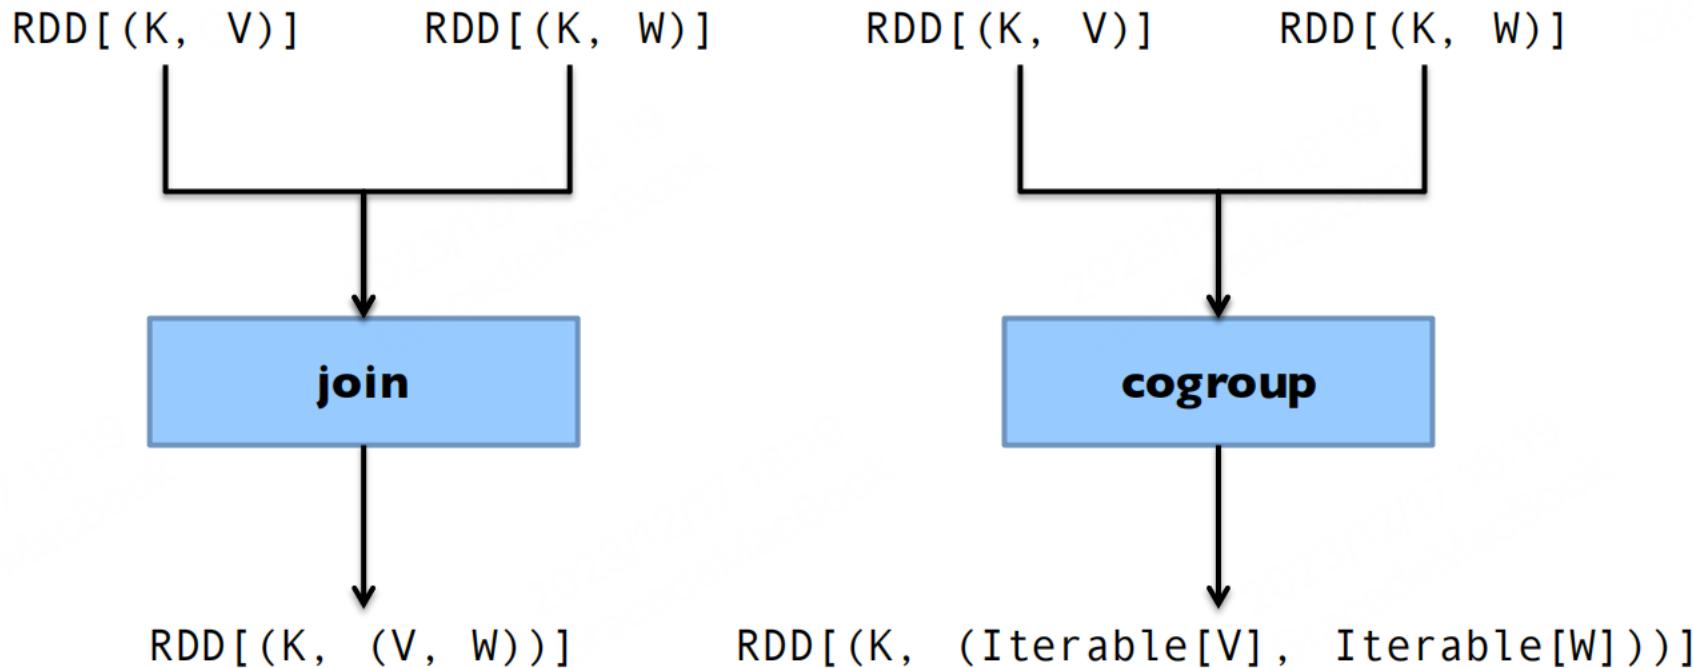
Not meant to be exhaustive

# Sort operations



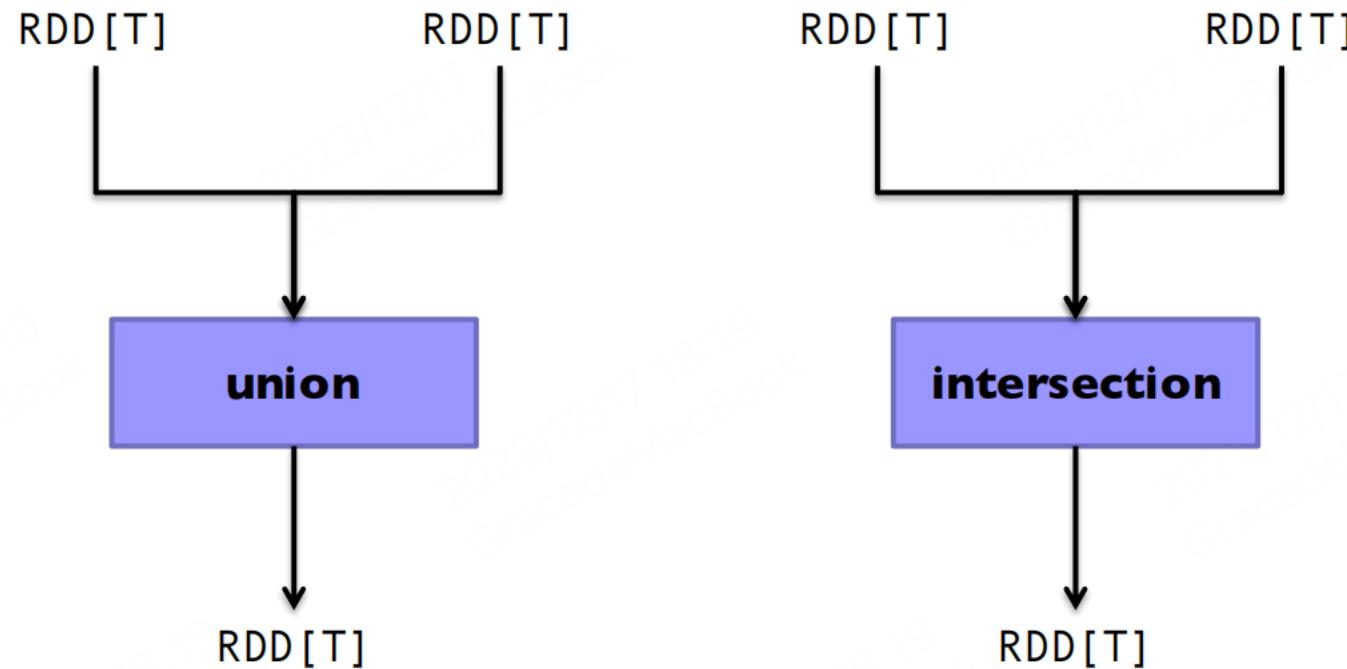
Not meant to be exhaustive

# Join-like operations



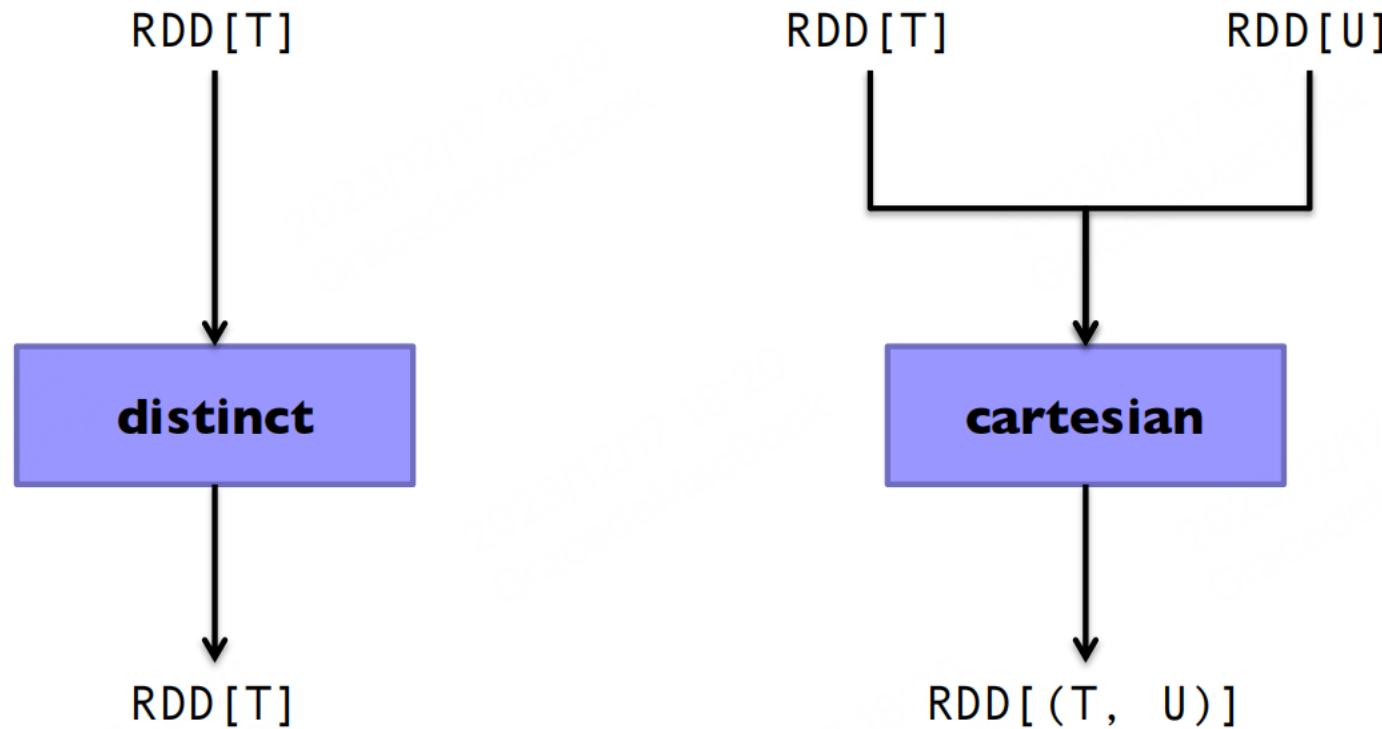
Not meant to be exhaustive

# Set-ish operations



Not meant to be exhaustive

# Set-ish operations



Not meant to be exhaustive

# Spark WordCount: 3-liner

```
text_file = sc.textFile("hdfs://...")  
counts = text_file.flatMap(lambda line: line.split(" ")) \  
    .map(lambda word: (word, 1)) \  
    .reduceByKey(lambda a, b: a + b)  
counts.saveAsTextFile("hdfs://...")
```

This?

```
private static class MyMapper  
    extends Mapper<LongWritable, Text, Text, IntWritable> {  
  
    private final static IntWritable ONE = new IntWritable(1);  
    private final static Text WORD = new Text();  
  
    @Override  
    public void map(LongWritable key, Text value, Context context)  
        throws IOException, InterruptedException {  
        String line = ((Text) value).toString();  
        String[] words = line.trim().split("\\s+");  
        for (String w: words) {  
            WORD.set(w);  
            context.write(WORD, ONE);  
        }  
    }  
  
    private static class MyReducer  
        extends Reducer<Text, IntWritable, Text, IntWritable> {  
  
        private final static IntWritable SUM = new IntWritable();  
  
        @Override  
        public void reduce(Text key, Iterable<IntWritable> values,  
                          Context context) throws IOException, InterruptedException {  
            int sum = 0;  
            for (IntWritable v: values) {  
                sum += v.get();  
            }  
            SUM.set(sum);  
            context.write(key, SUM);  
        }  
  
        Configuration conf = getConf();  
        Job job = Job.getInstance(conf);  
        job.setJobName("Word Count");  
        job.setJarByClass(WordCount.class);  
        job.setNumReduceTasks(reduceTasks);  
        FileInputFormat.setInputPaths(job, new Path(inputPath));  
        FileOutputFormat.setOutputPath(job, new Path(outputPath));  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        job.setMapperClass(WordCountMapper.class);  
        job.setCombinerClass(WordCountReducer.class);  
        job.setReducerClass(WordCountReducer.class);  
        System.exit(job.waitForCompletion(true) ? 0: 1);  
    }  
}
```

Or this?

# Credits

- Some slides are adapted from course slides of COMP 4651 in HKUST
- Some slides adapted from course slides of DS5110 in UVA